

CSCI/ECEN 3302

Introduction to Robotics

Alessandro Roncone

aroncone@colorado.edu

<https://hiro-group.ronc.one>

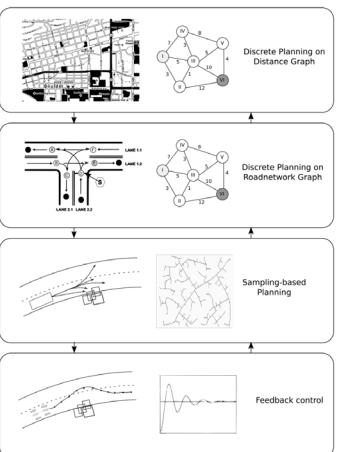
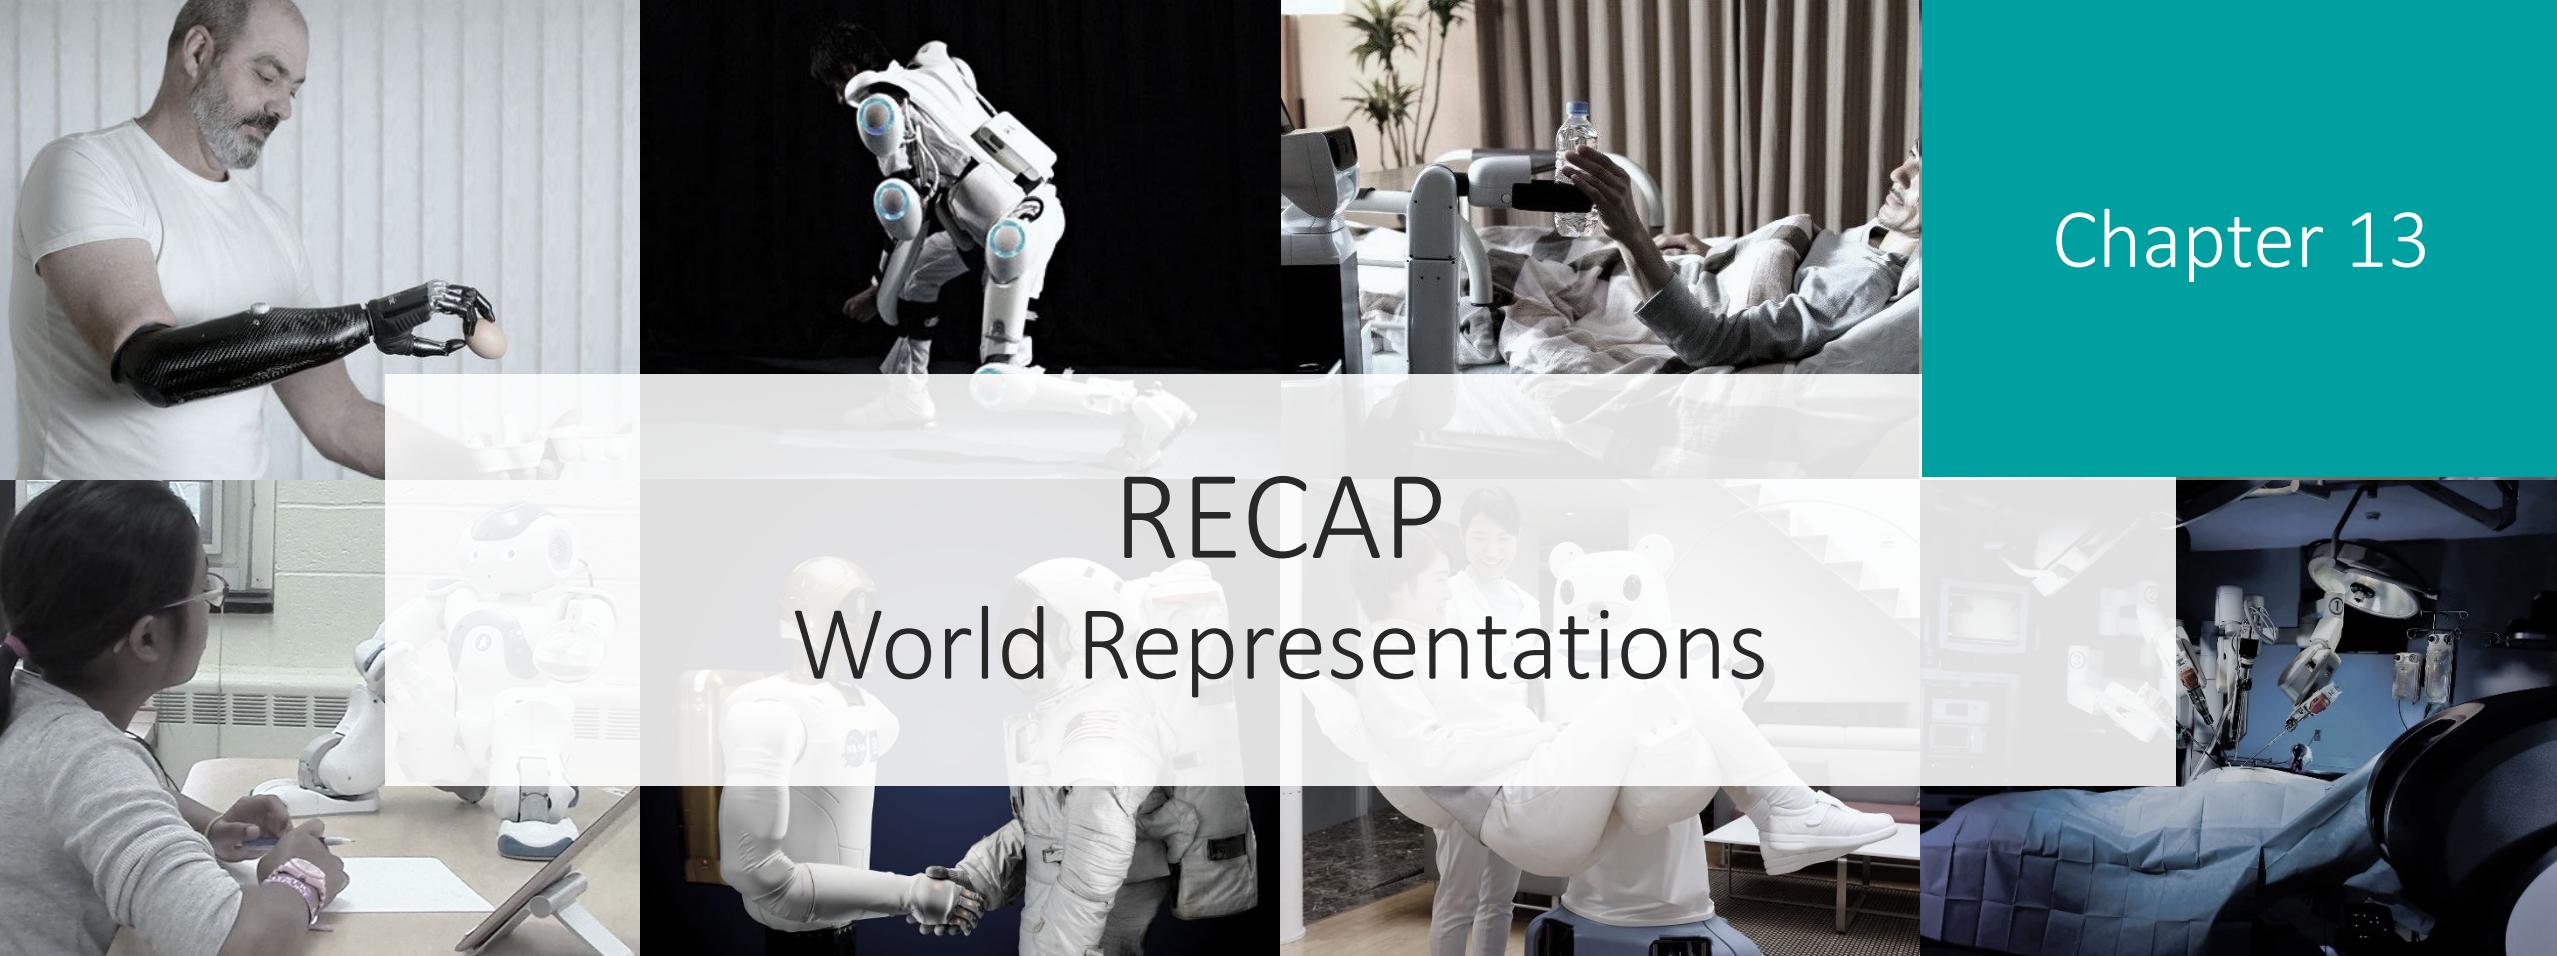
Administrivia

- Lab 2
 - Grades are out
 - Next in line: HW1, then Lab 3, then Lab 4/5
- Lab 4
 - Deadline **tonight!**
- Lab 5
 - Three weeks long
 - To be concluded before the Spring Break
- HW2 to come up soon – depending on how much I can cover today vs next Tue

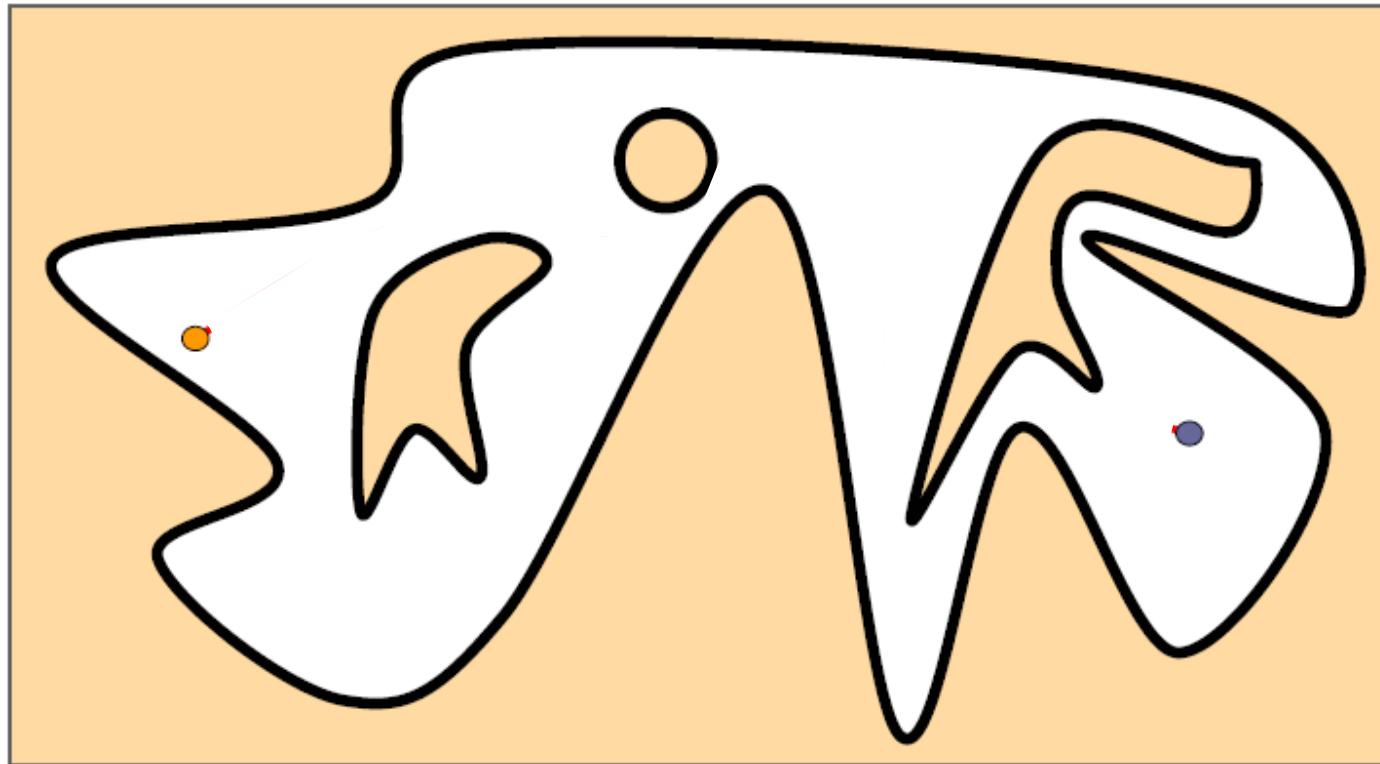
Chapter 13

RECAP

World Representations



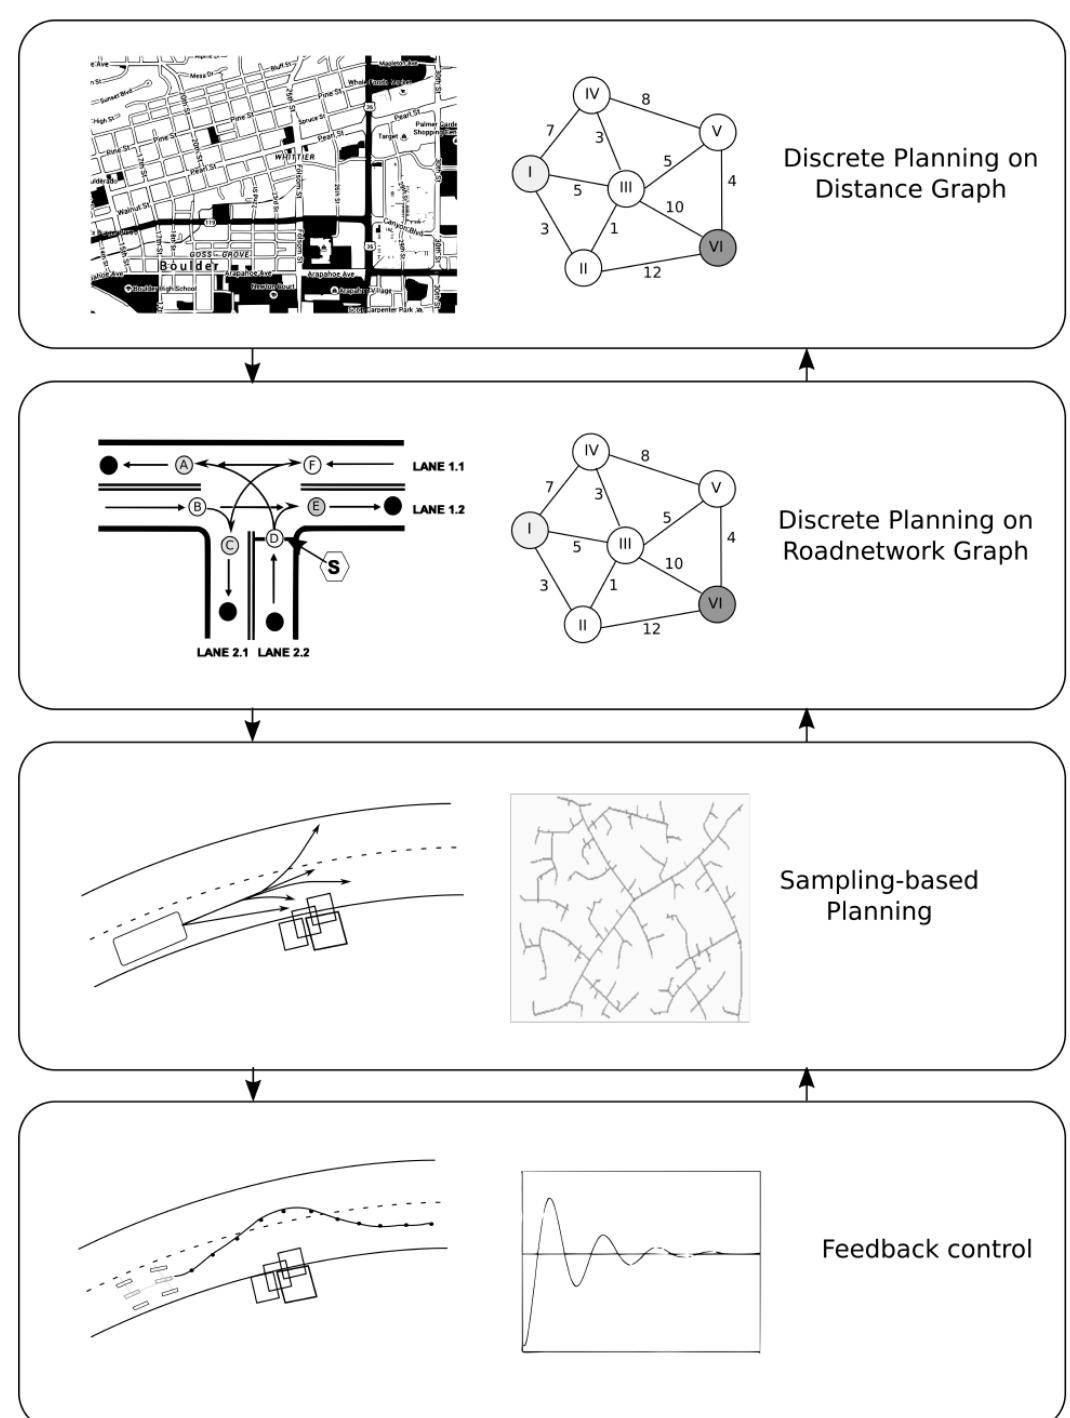
Moving from Start to Goal state



- General **approach**: Reduce what amounts to an intractable problem in continuous C-space to a tractable problem in a discrete space.
- **Tools**:
 - Environment Representations
 - Search Algorithms

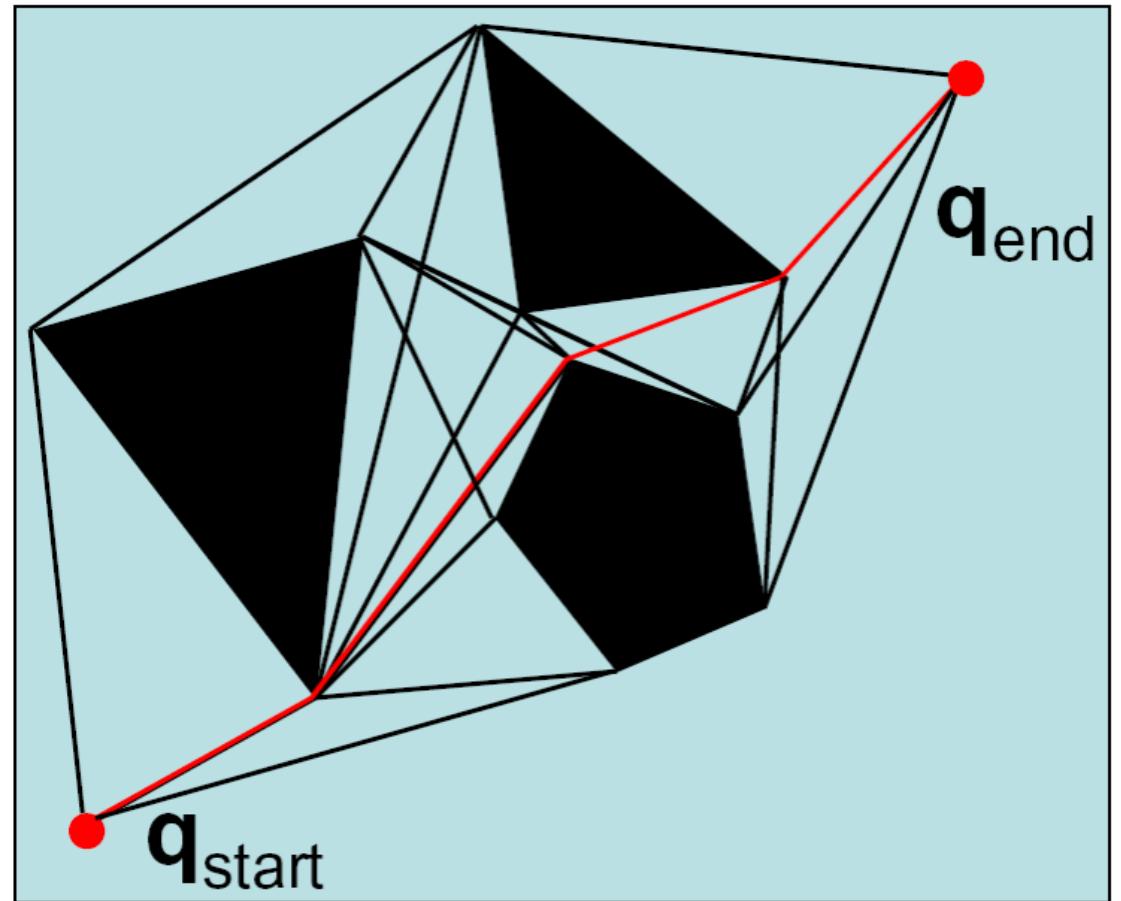


Planning across length scales



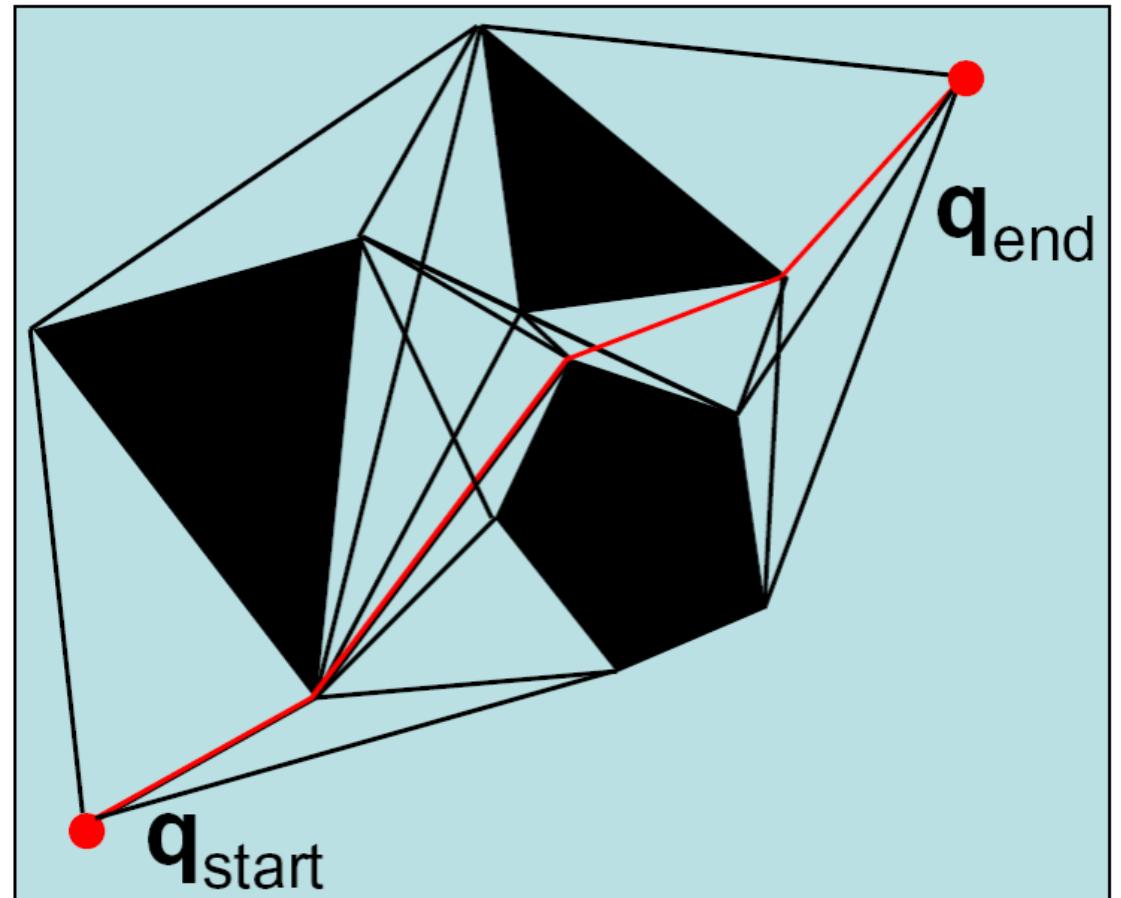
Some Well-Known Representations

- Visibility Graphs
- Roadmap
- Cell Decomposition
- Potential Field



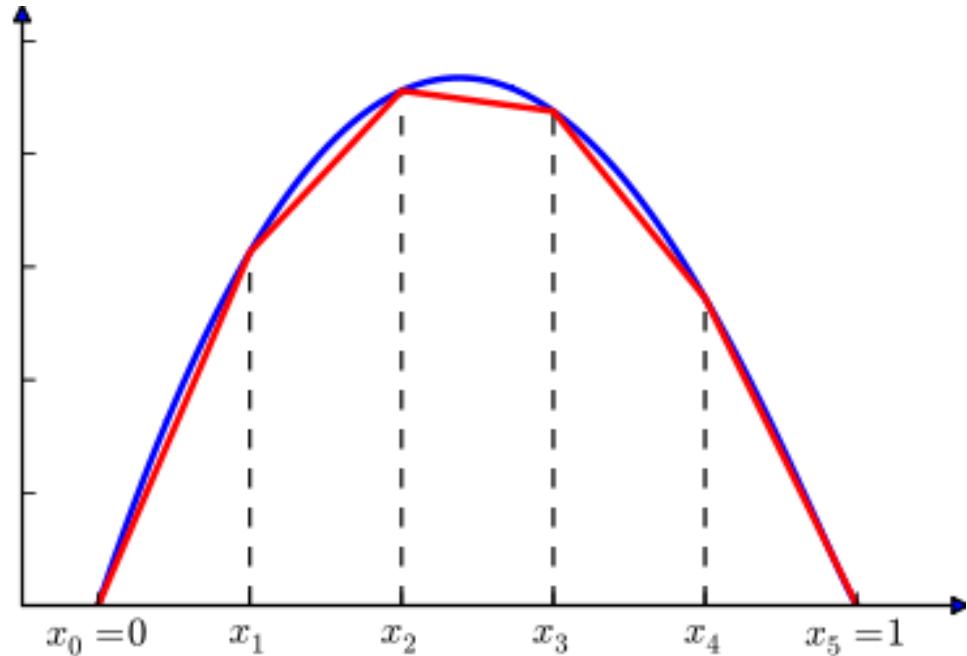
Some Well-Known Representations

- Visibility Graphs
- Roadmap
- Cell Decomposition
- Potential Field



Visibility Graph Method

- If there is a collision-free path between two points, then there is a polygonal path that bends only at the obstacles vertices.
- A polygonal path is a piecewise linear curve:

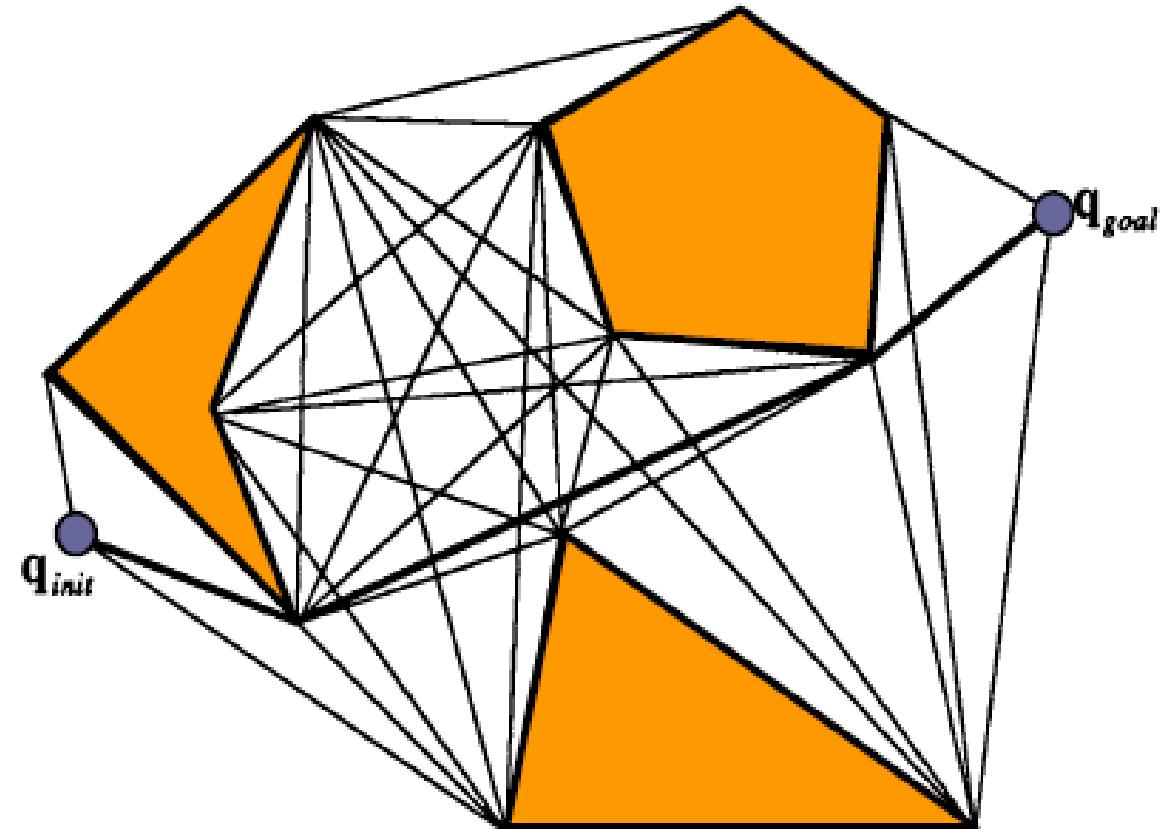


Visibility Graph Algorithm

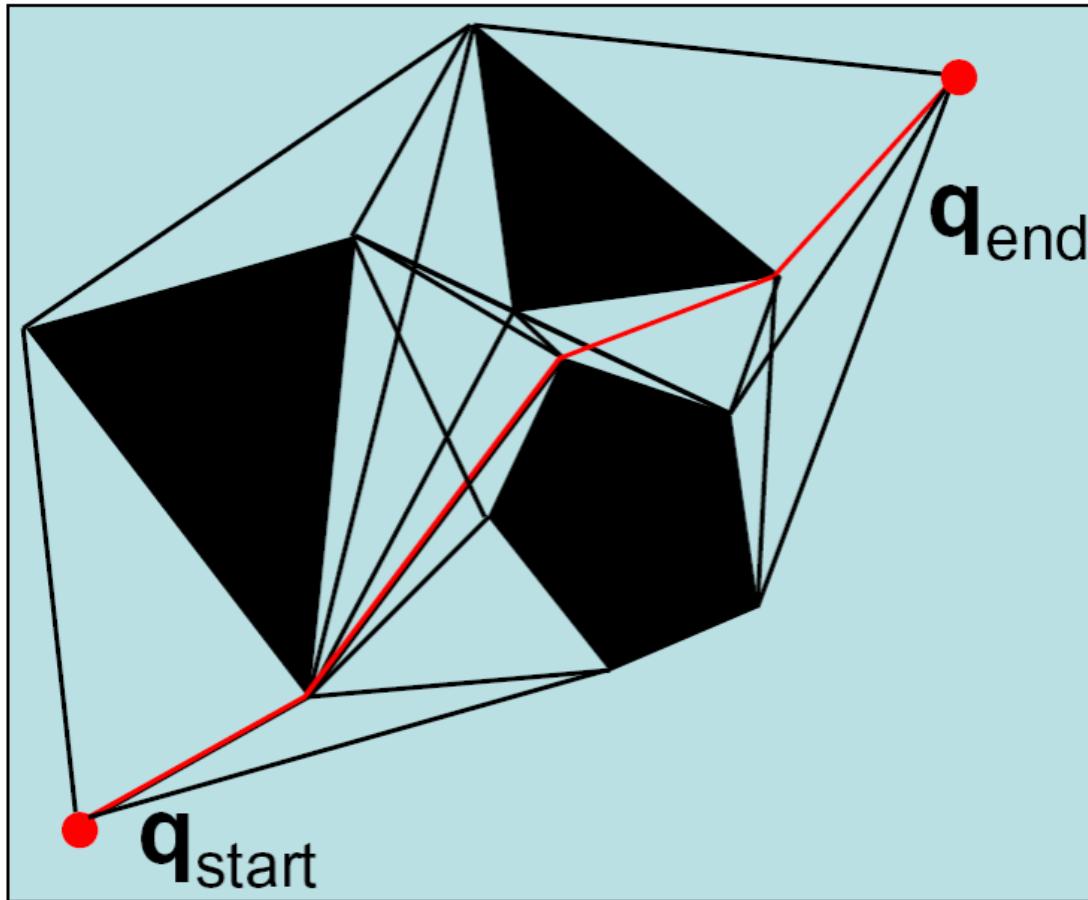
Input: q_{init} , q_{goal} , polygonal obstacles

Output: visibility graph G

1. **for** every pair of nodes u, v
2. **if** segment (u, v) is an obstacle edge **then**
3. insert edge (u, v) into G;
4. **else**
5. **for** every obstacle edge e
6. **if** segment (u, v) intersects e
7. go to (1);
8. insert edge (u, v) into G.

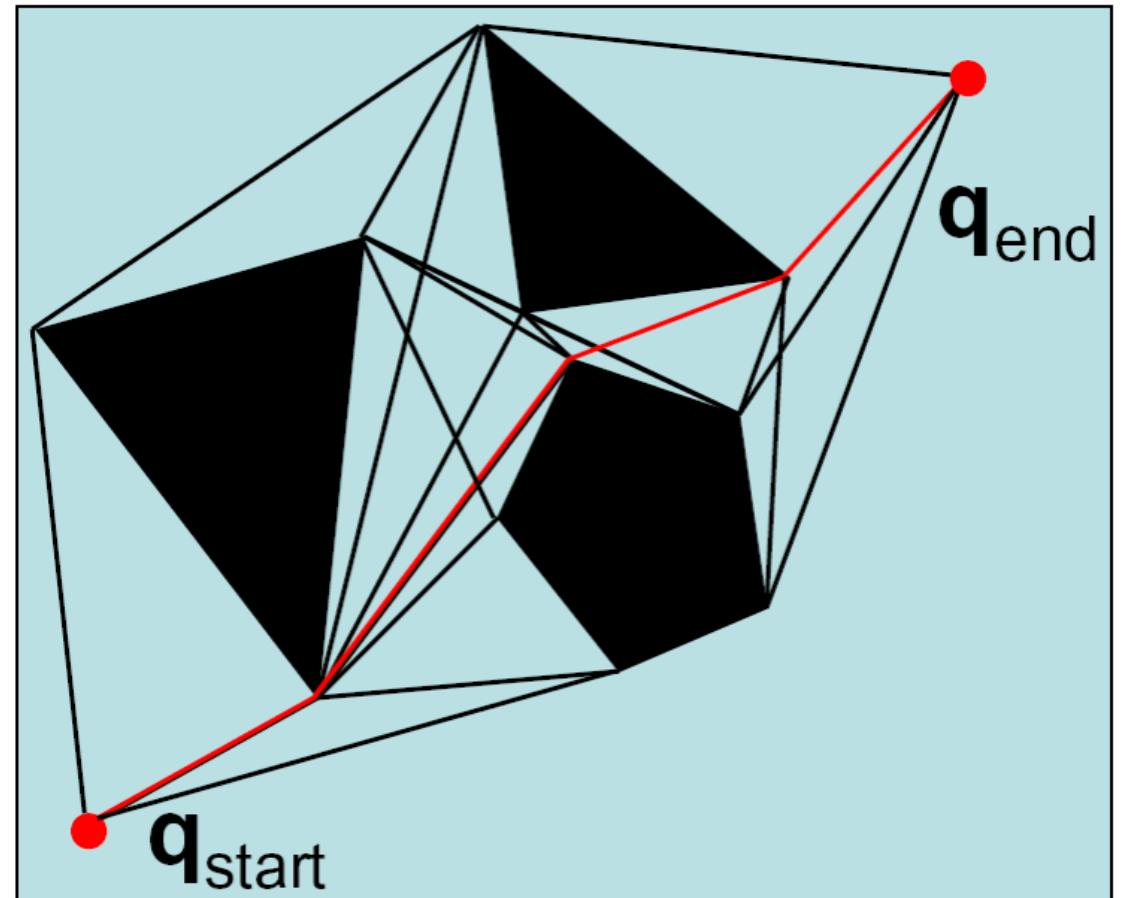


Visibility Graph: Strengths/Weaknesses?



Some Well-Known Representations

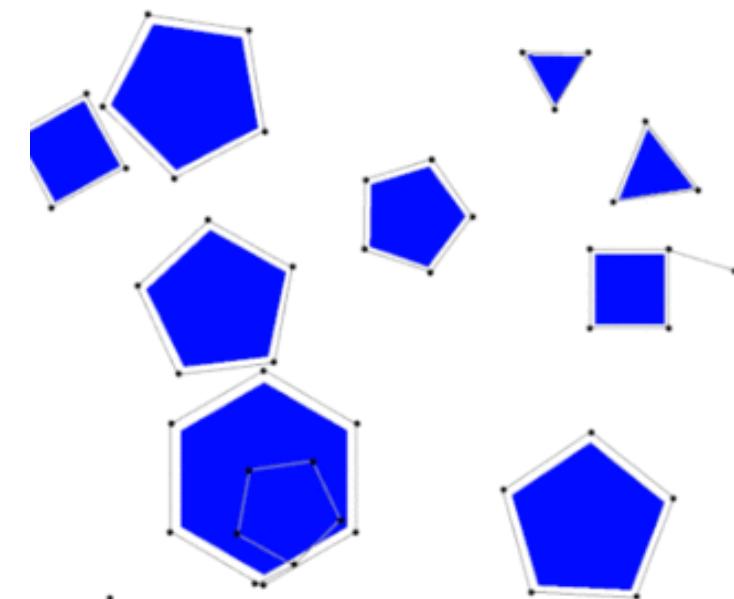
- Visibility Graphs
- Roadmap
- Cell Decomposition
- Potential Field



Road mapping

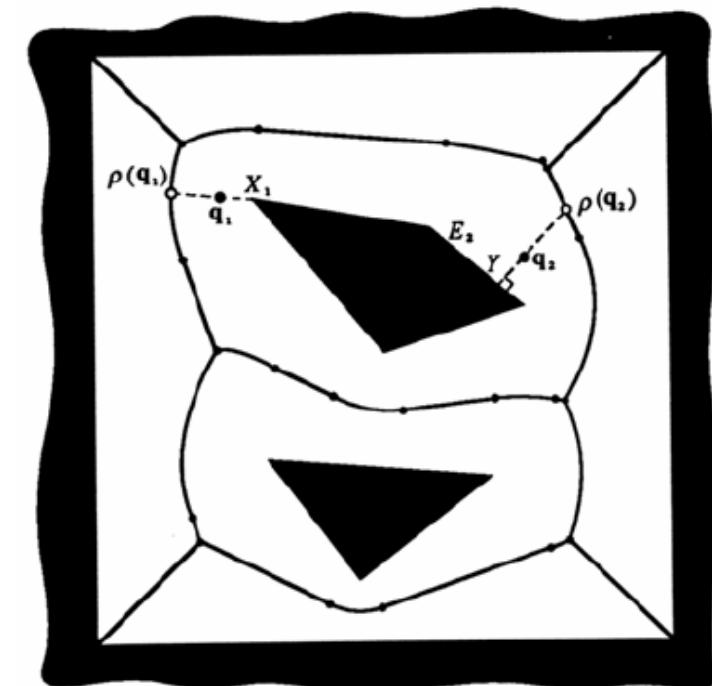
Probabilistic Road Mapping

- Take random samples, test them to see they are in free space, and connect them via nearest neighbor
- Use a graph search algorithm to find the optimal path



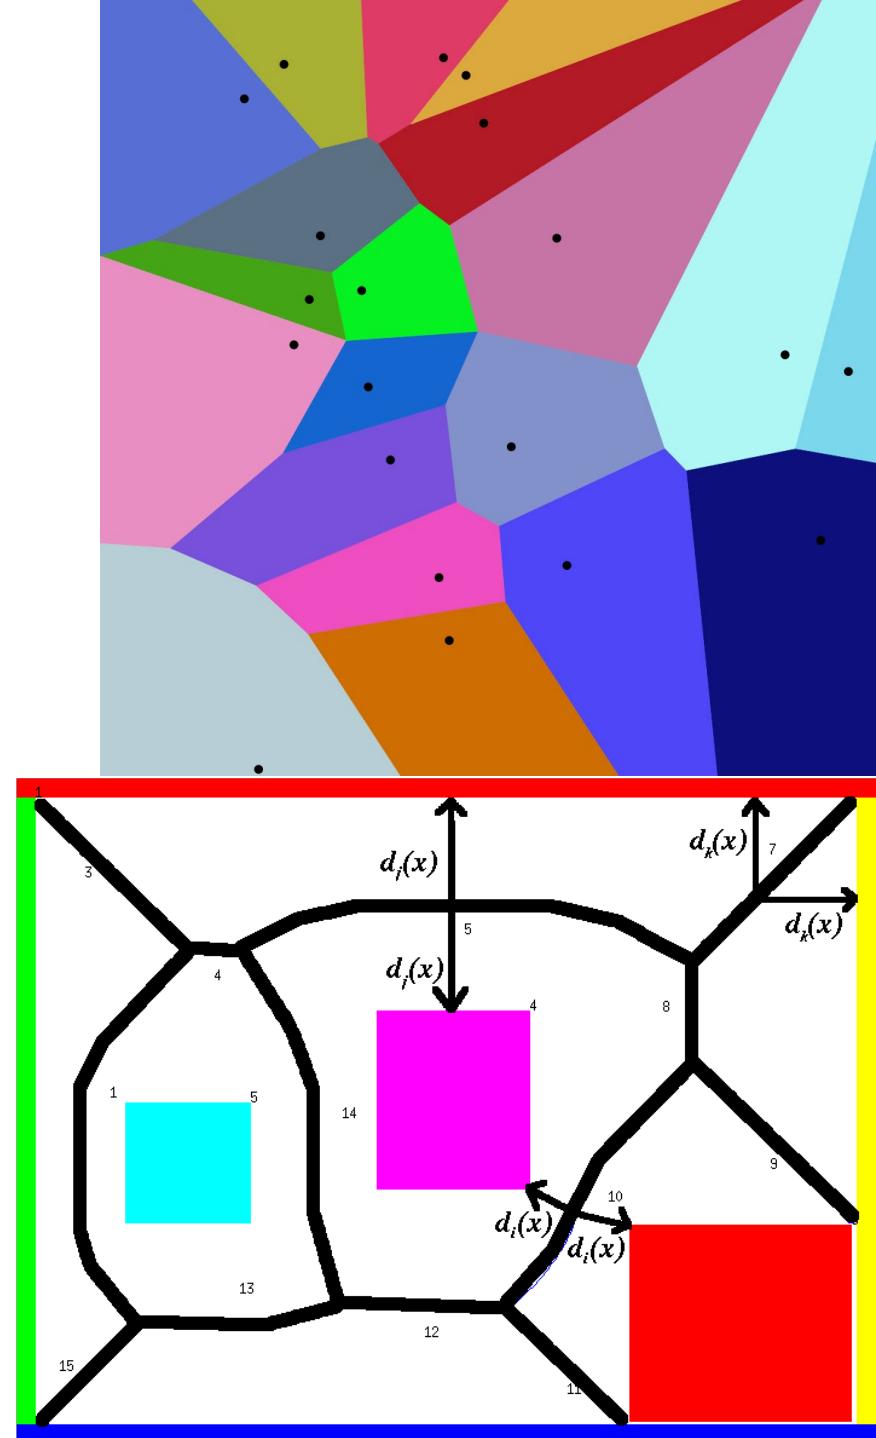
Voronoi Road Mapping

- Introduced by computational geometry researchers.
- Generate paths that maximize clearance
- Applicable mostly to 2-D configuration spaces

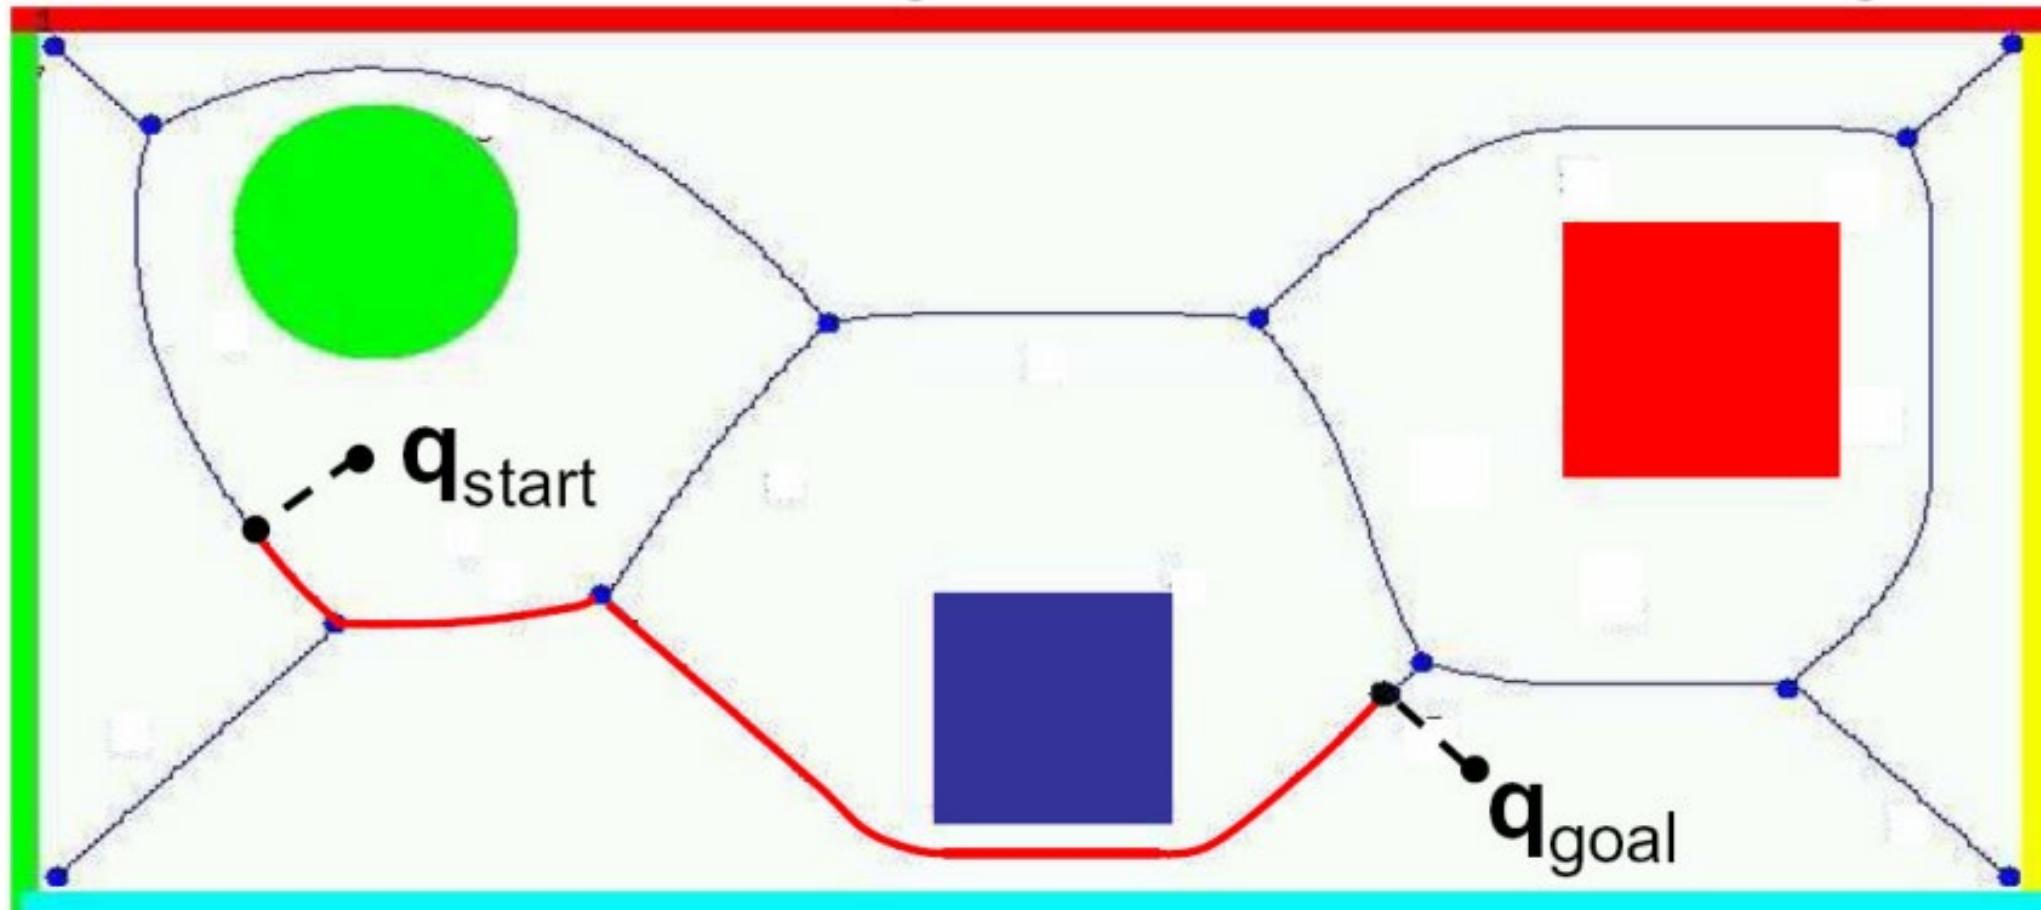


Voronoi Road Mapping

- Voronoi Diagram: partition of a plan with n points into n convex polygons whose points are closer to the generating point than any other
 - It is based on the topology of the environment
 - It creates a skeleton-like structure obtained via retraction, i.e. by progressively shrinking the free space.
 - Solutions have maximum clearance!

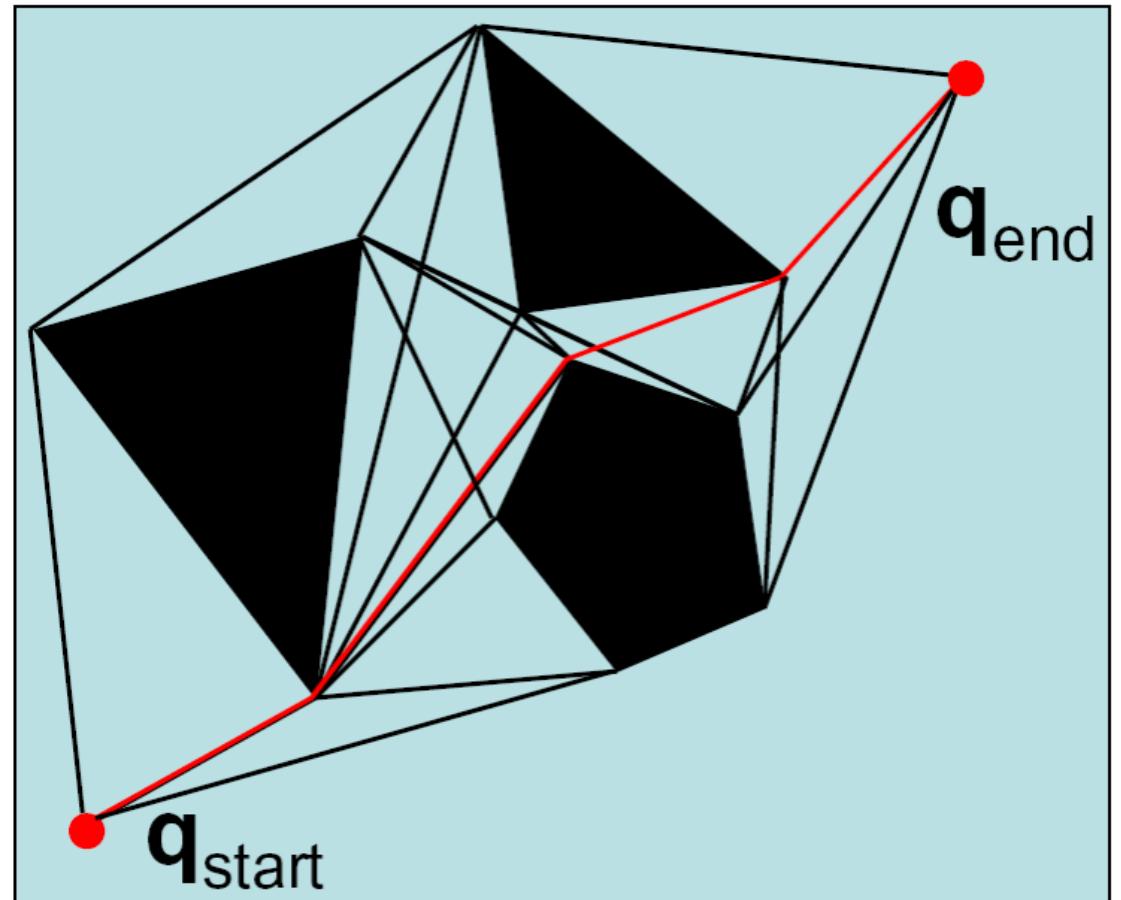


Voronoi Road Mapping: Strengths / Weaknesses

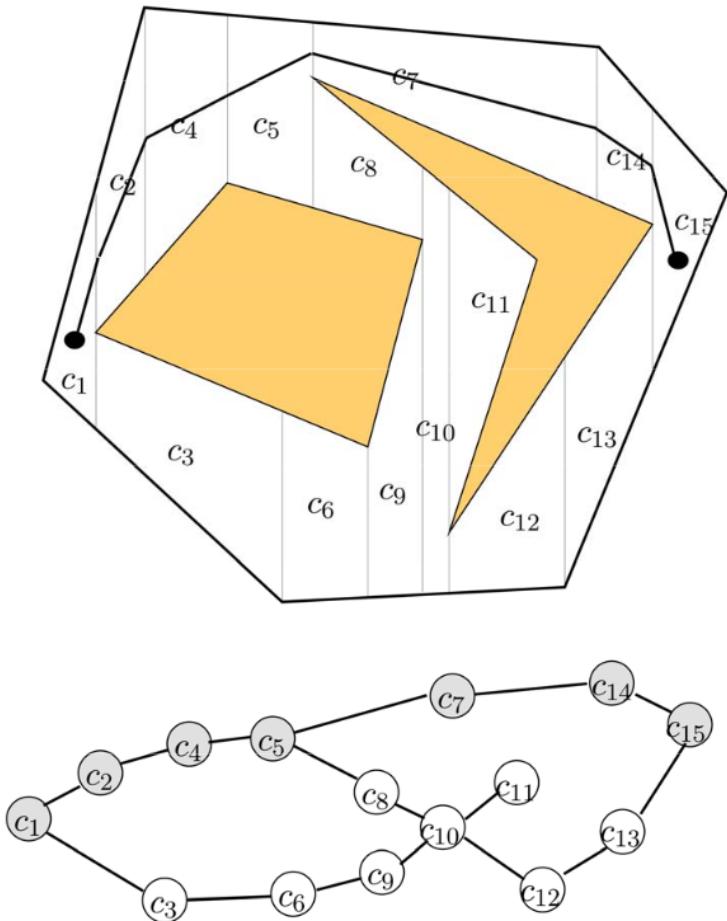


Some Well-Known Representations

- Visibility Graphs
- Roadmap
- Cell Decomposition
- Potential Field



Cell Decomposition



Exact cell decomposition

- Divides a space F precisely into sub-units

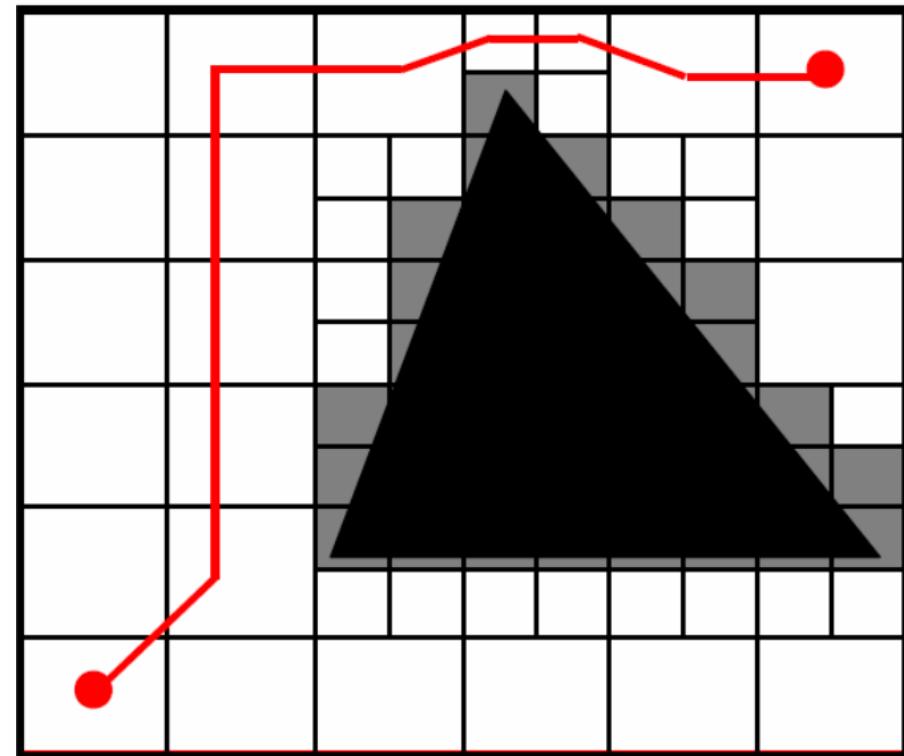
Approximate cell decomposition

- F is represented by a collection of non-overlapping cells whose union is contained in F .
- Cells usually have simple, regular shapes, e.g., rectangles, squares.
- Facilitates hierarchical space decomposition

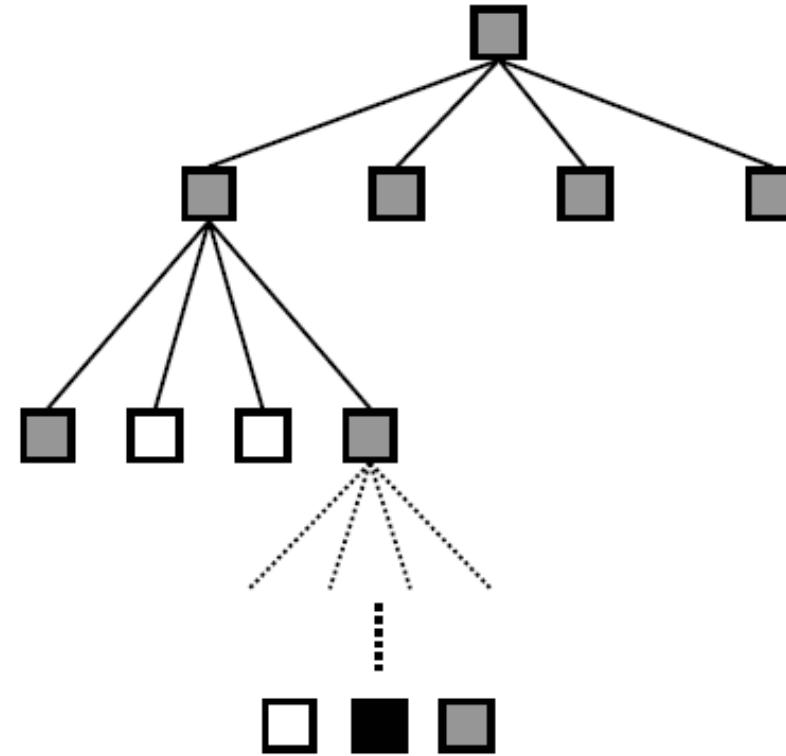
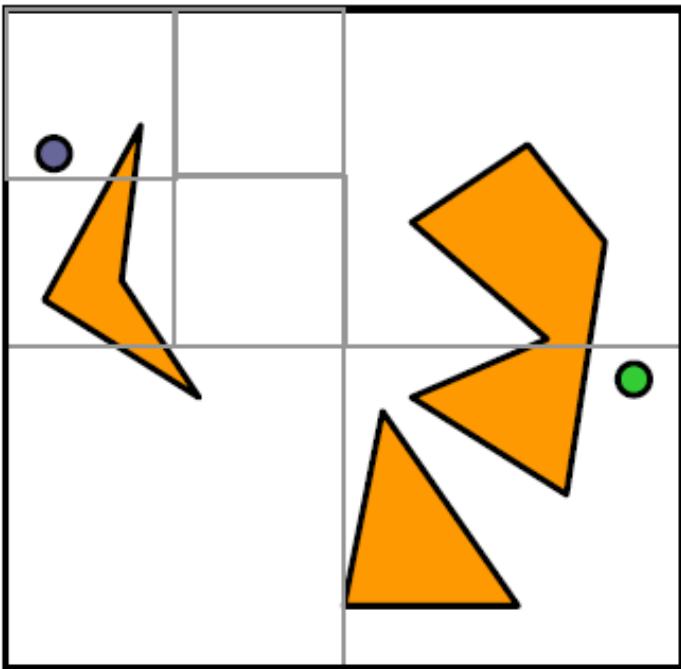
Cell Decomposition

Not necessarily *complete*

(Complete: If a solution exists, it will eventually be found)



Quadtree Decomposition

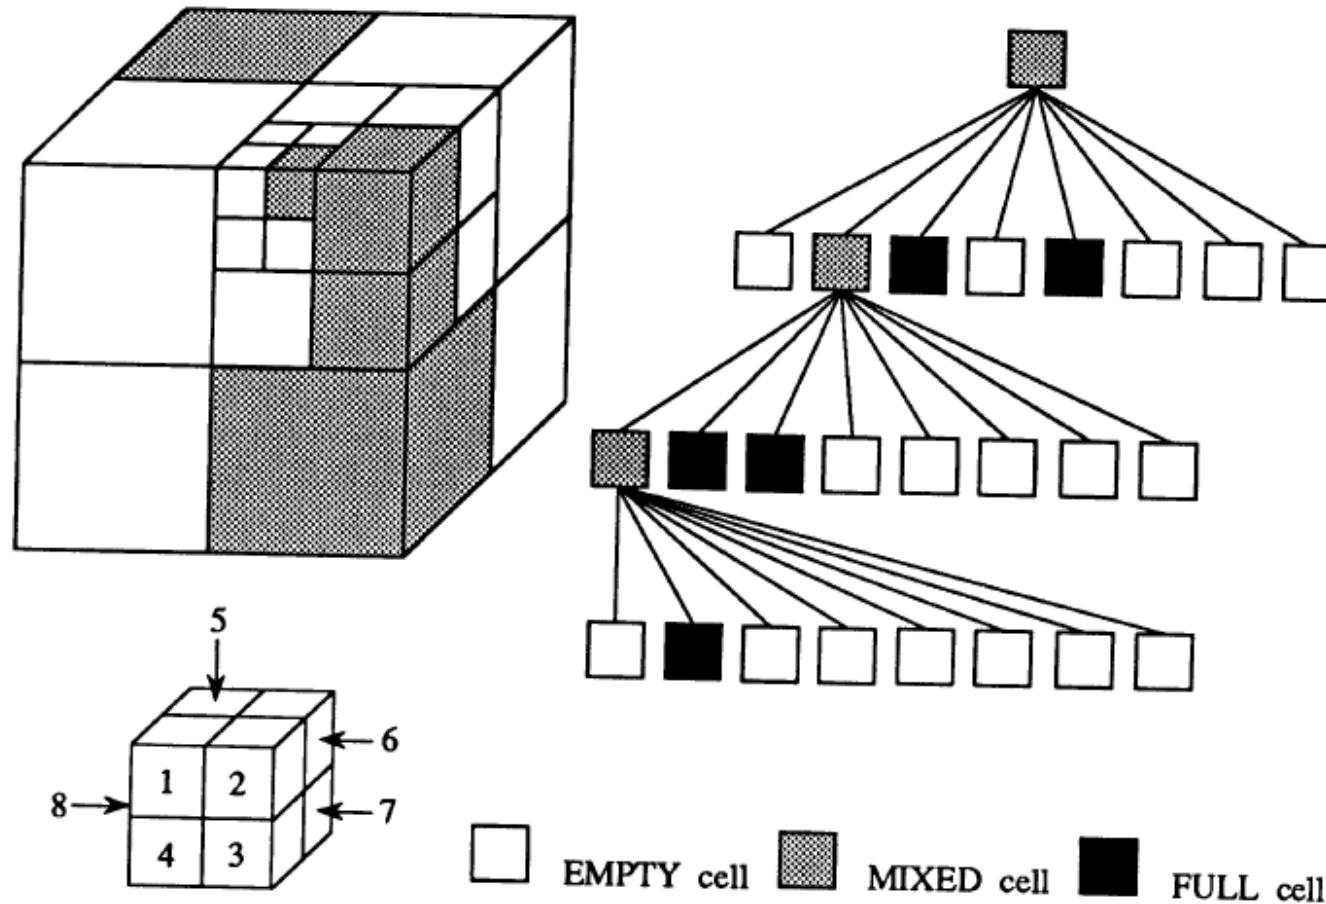


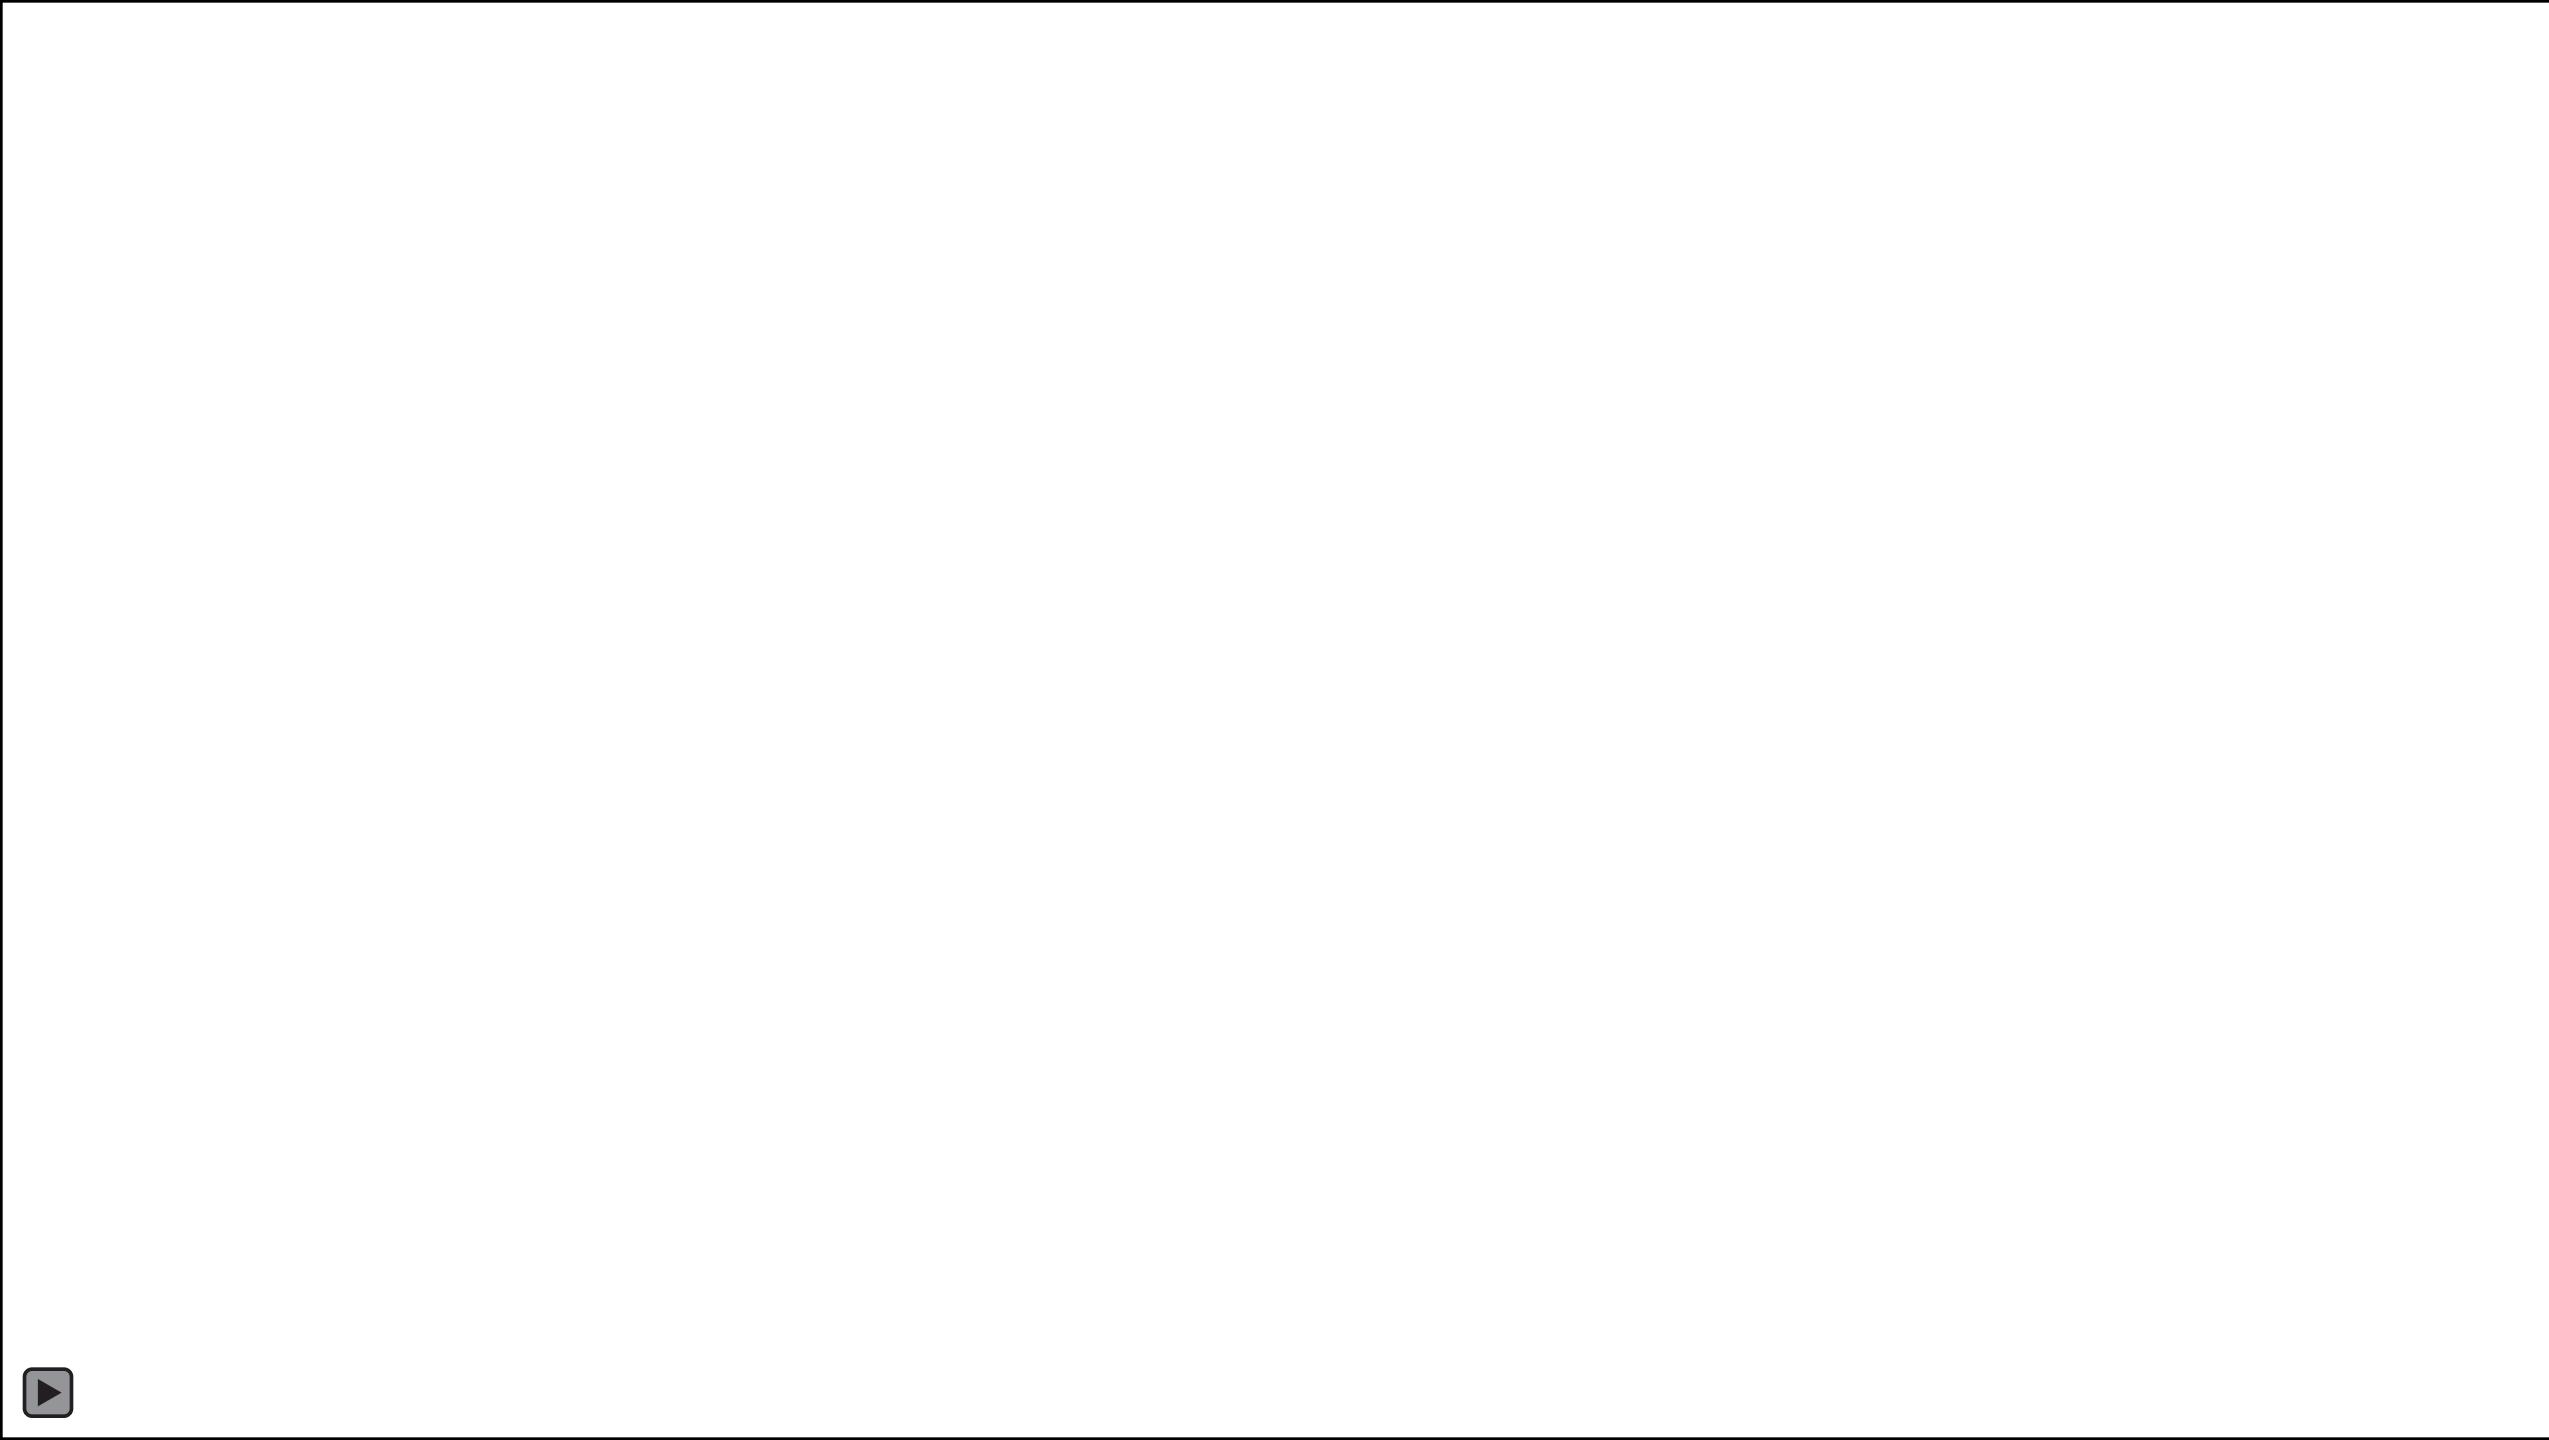
□ empty

■ mixed

■ full

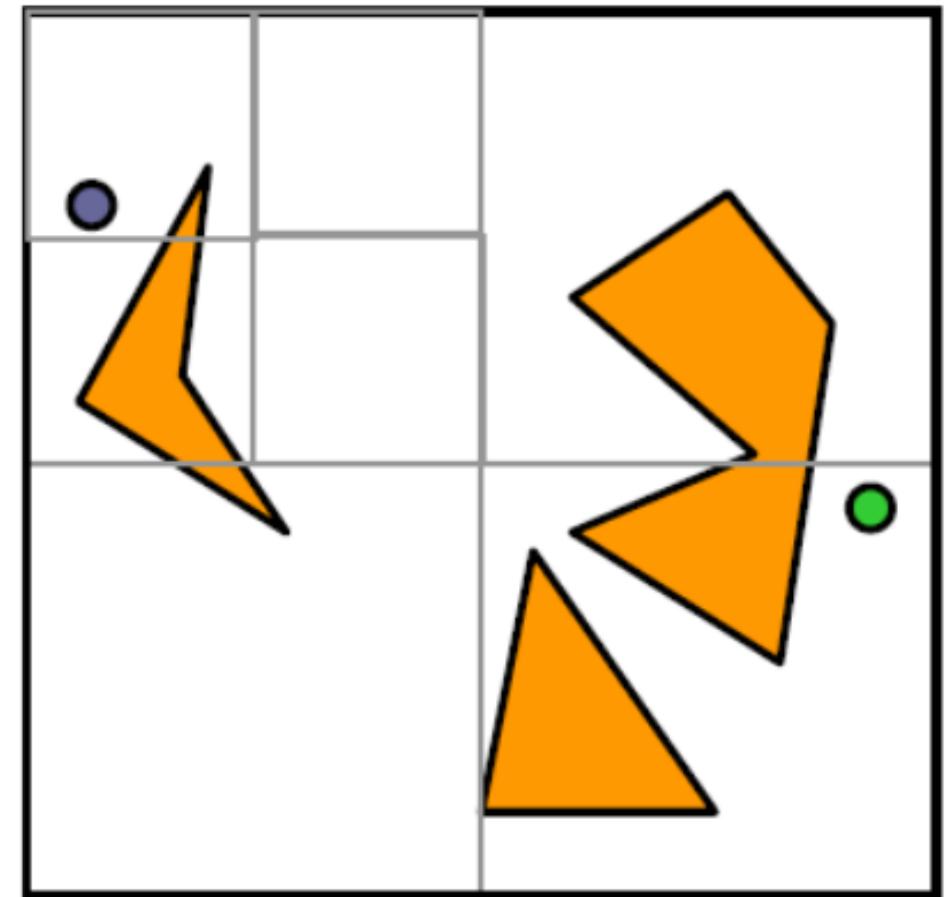
Octree Decomposition



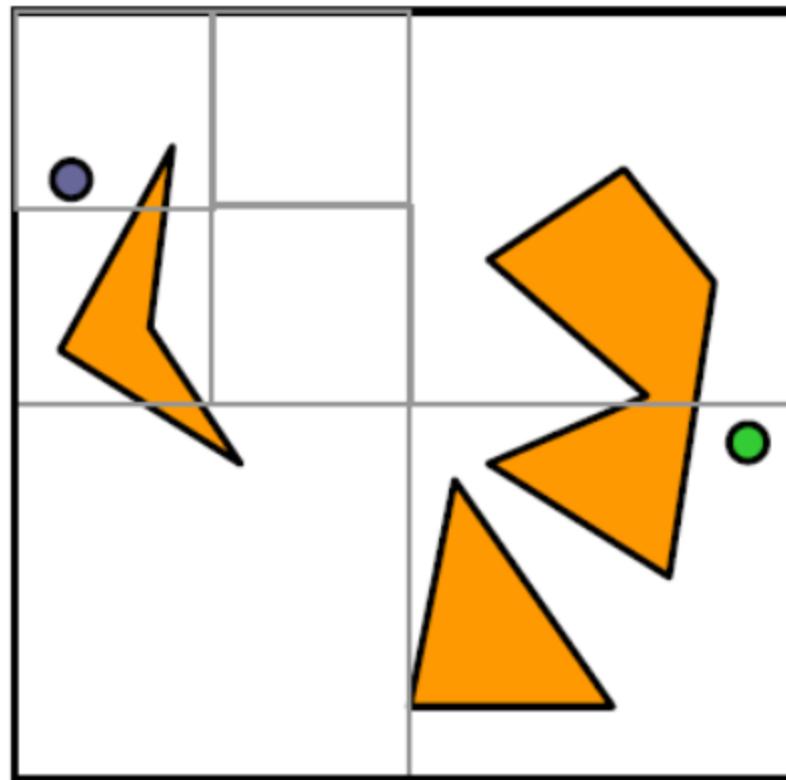


Cell Decomposition: Algorithm Outline

1. Decompose the free space F into cells
2. Search for a sequence of mixed or free cells that connect the initial and goal positions
3. Further decompose the mixed
4. Repeat (2) and (3) until a sequence of free cells is found

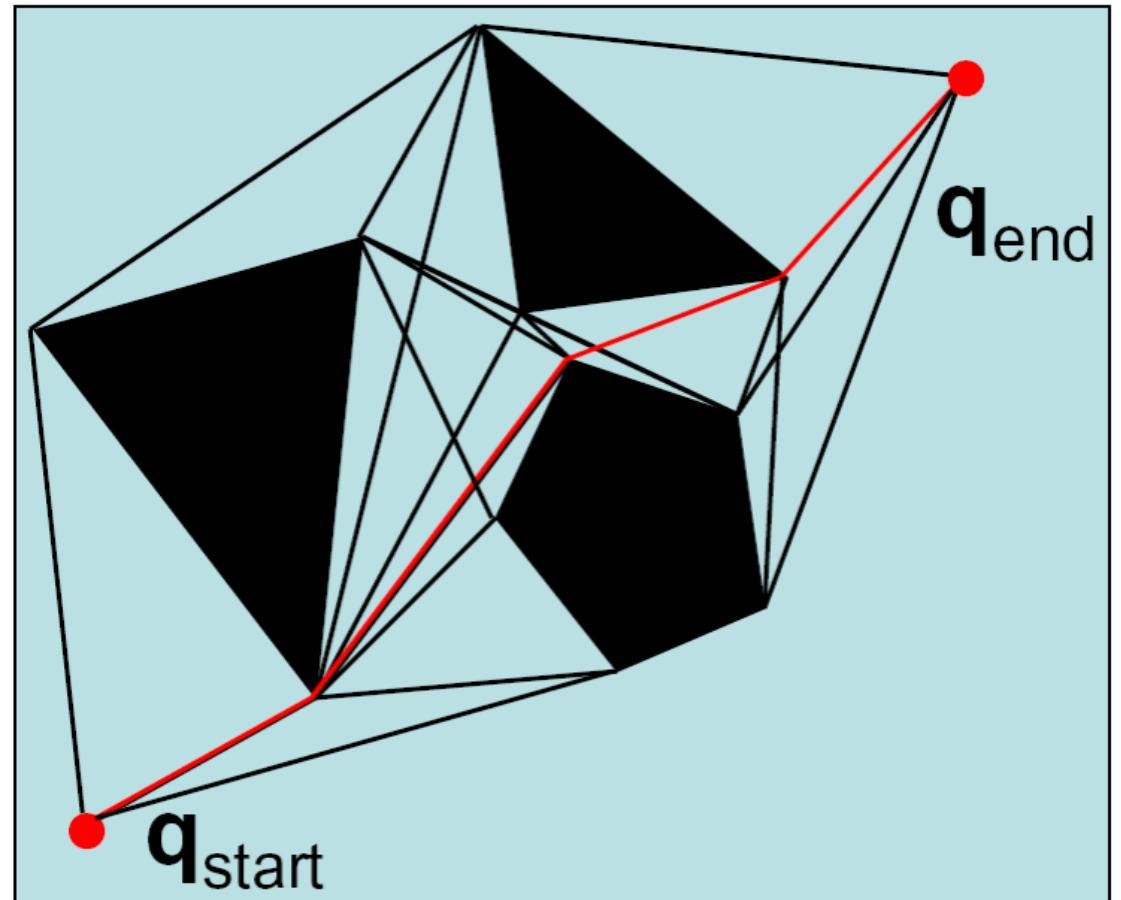


Cell Decomposition: Strengths/Weaknesses



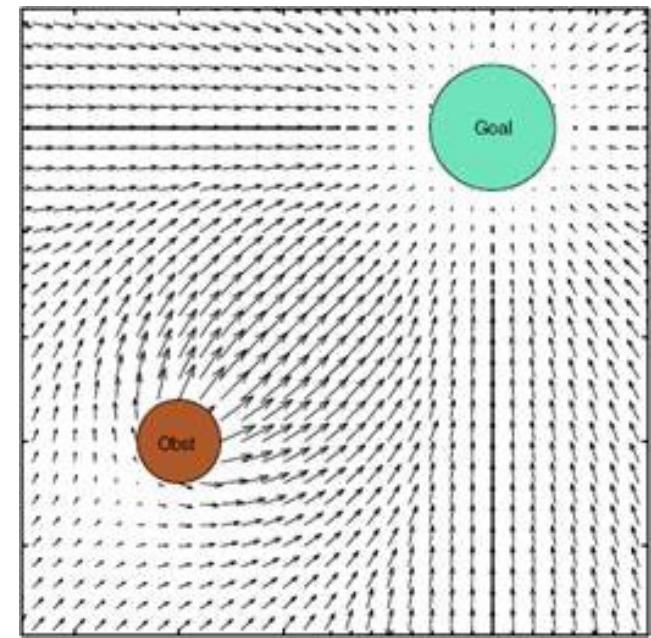
Some Well-Known Representations

- Visibility Graphs
- Roadmap
- Cell Decomposition
- Potential Field

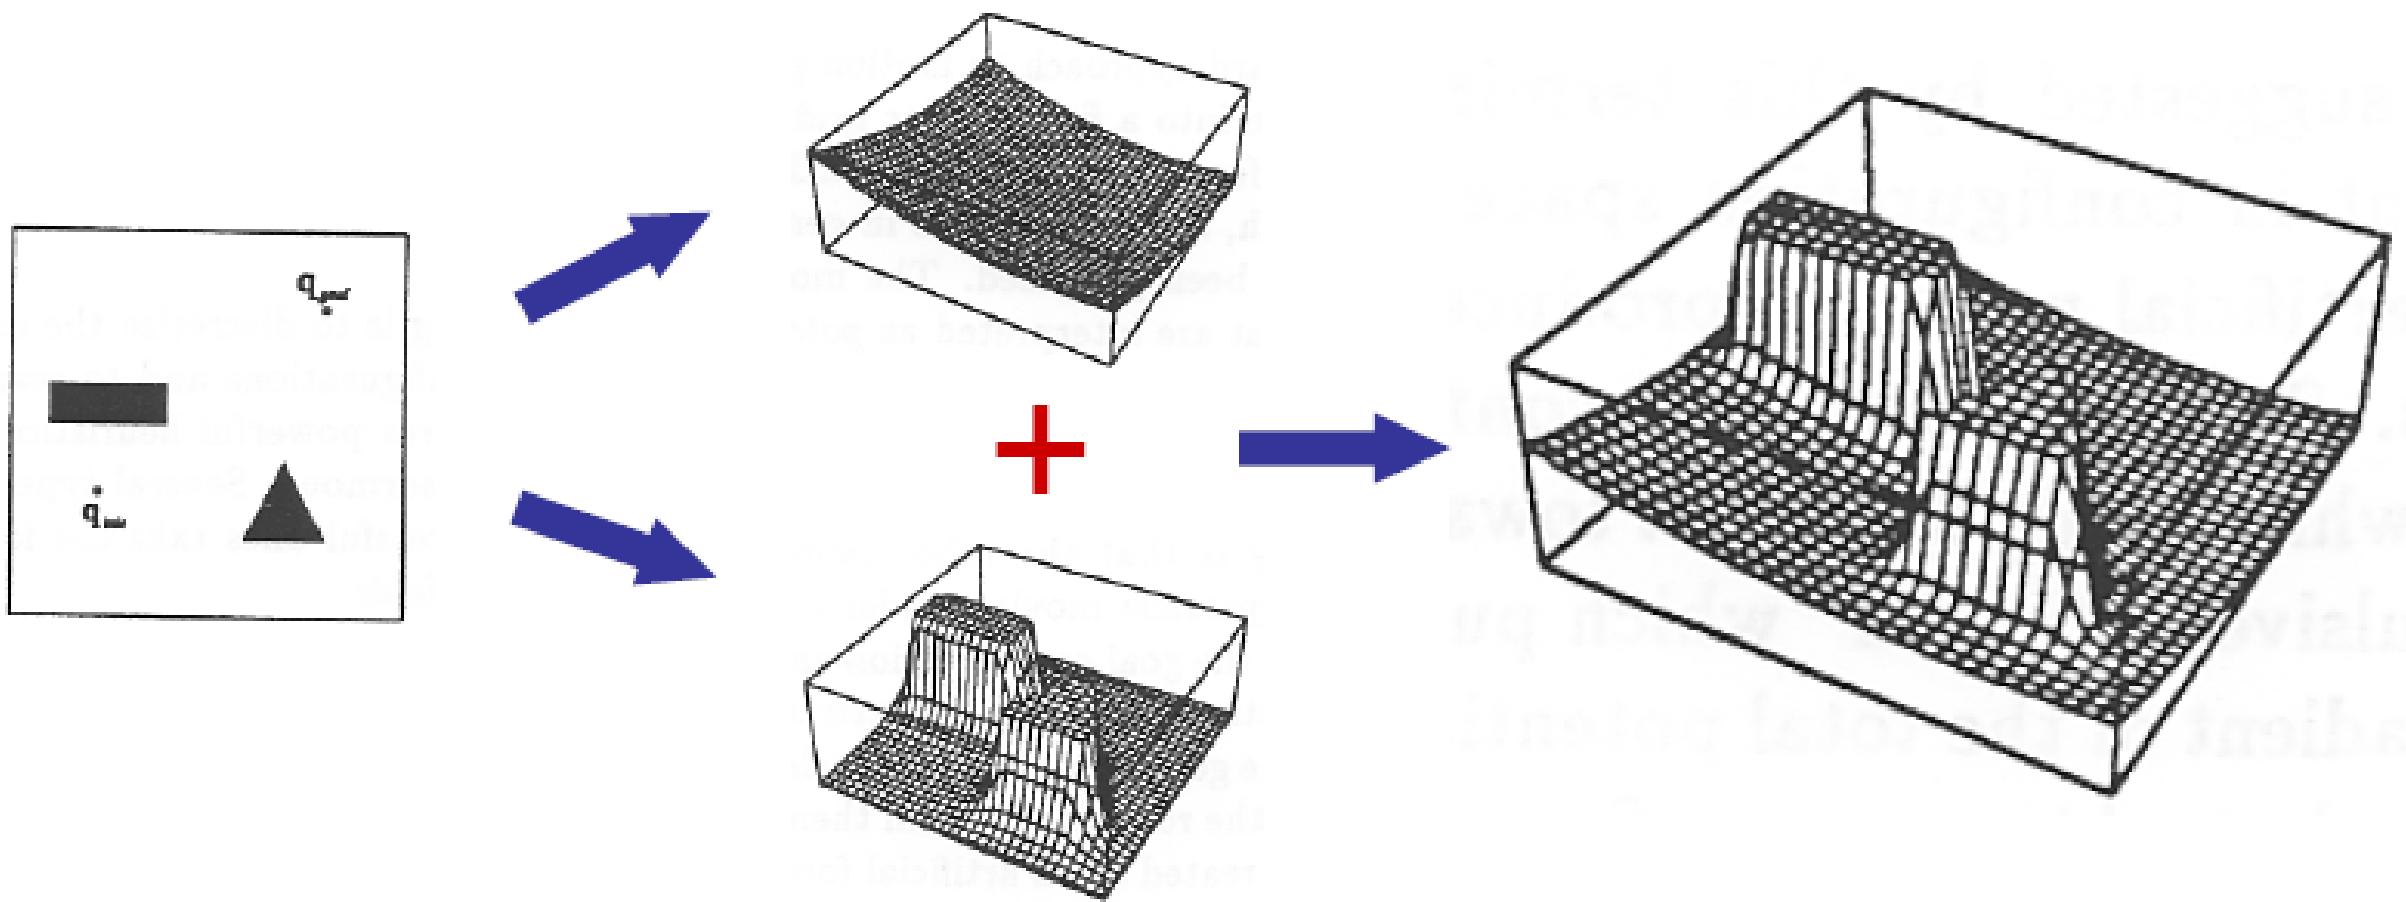


Potential Fields

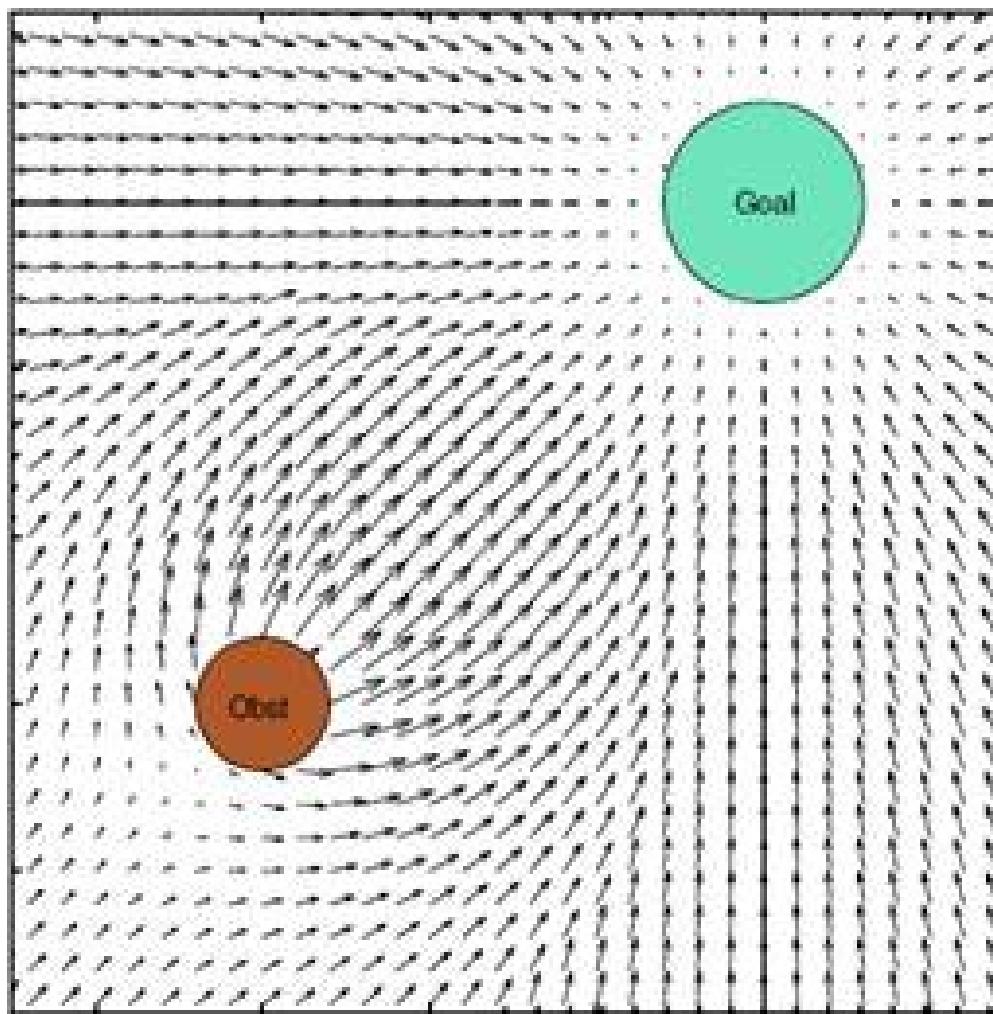
- Initially proposed for real-time collision avoidance [Khatib 1986].
- A potential field is a scalar function over the free space.
- To navigate, the robot applies a force proportional to the negated gradient of the potential field.
- A navigation function is an ideal potential field that
 - has global minimum at the goal
 - has no local minima
 - grows to infinity near obstacles
 - is smooth



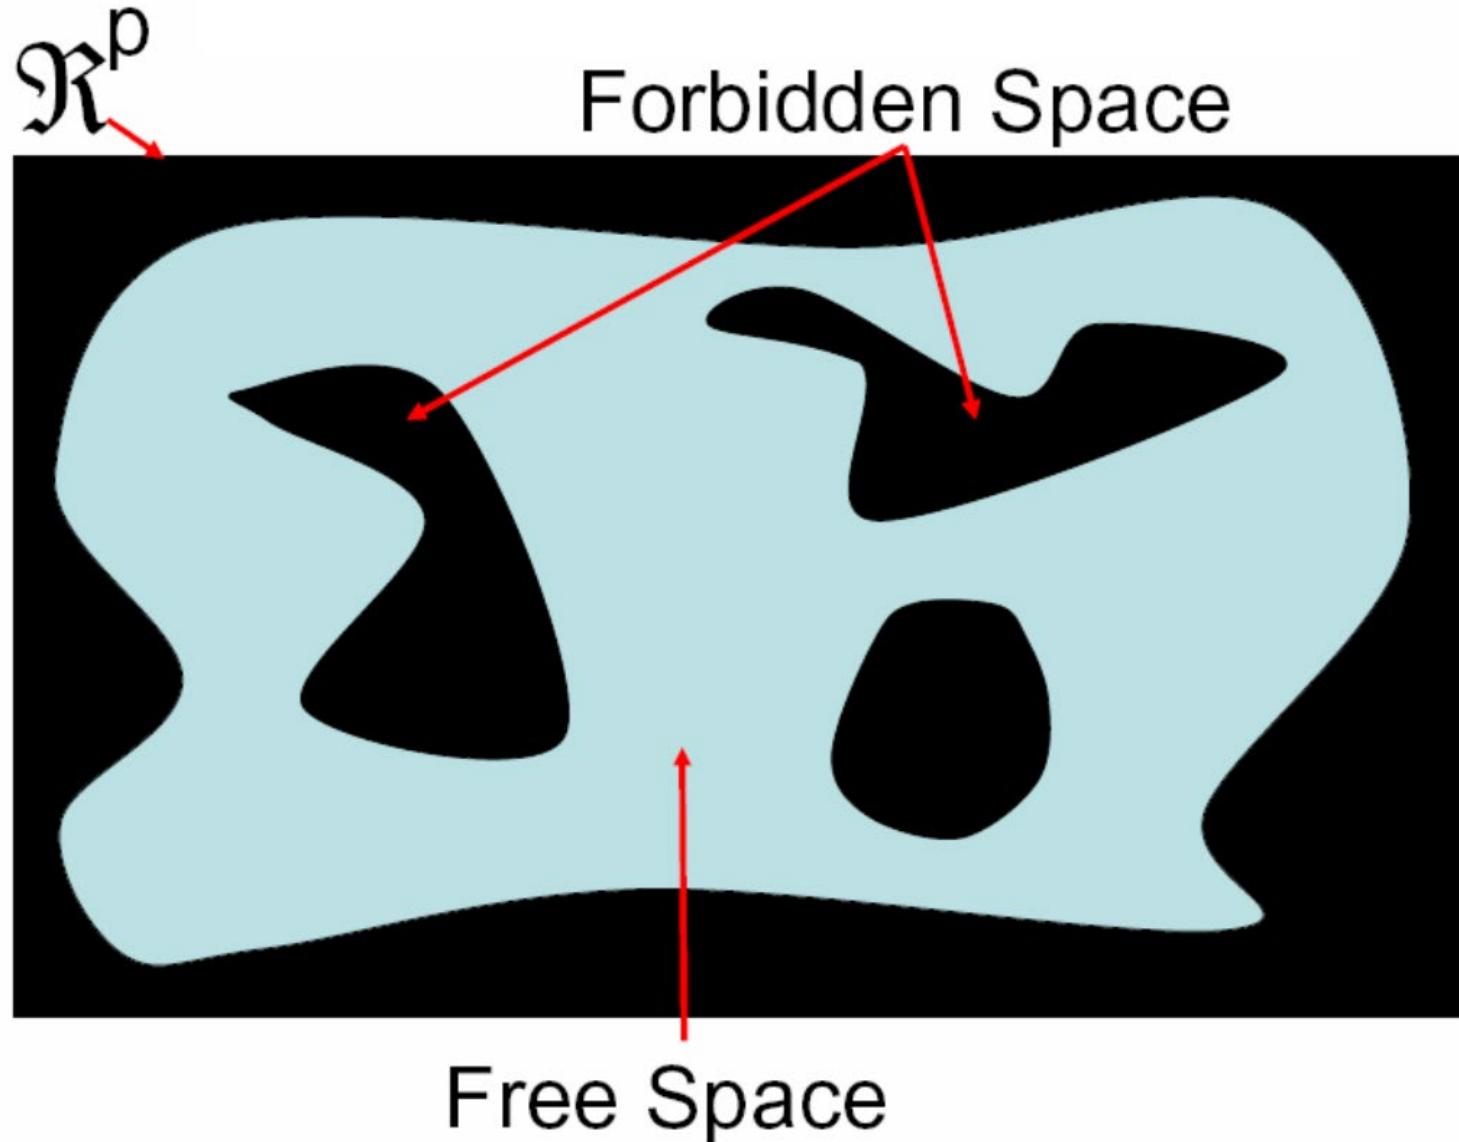
How Does It Work?



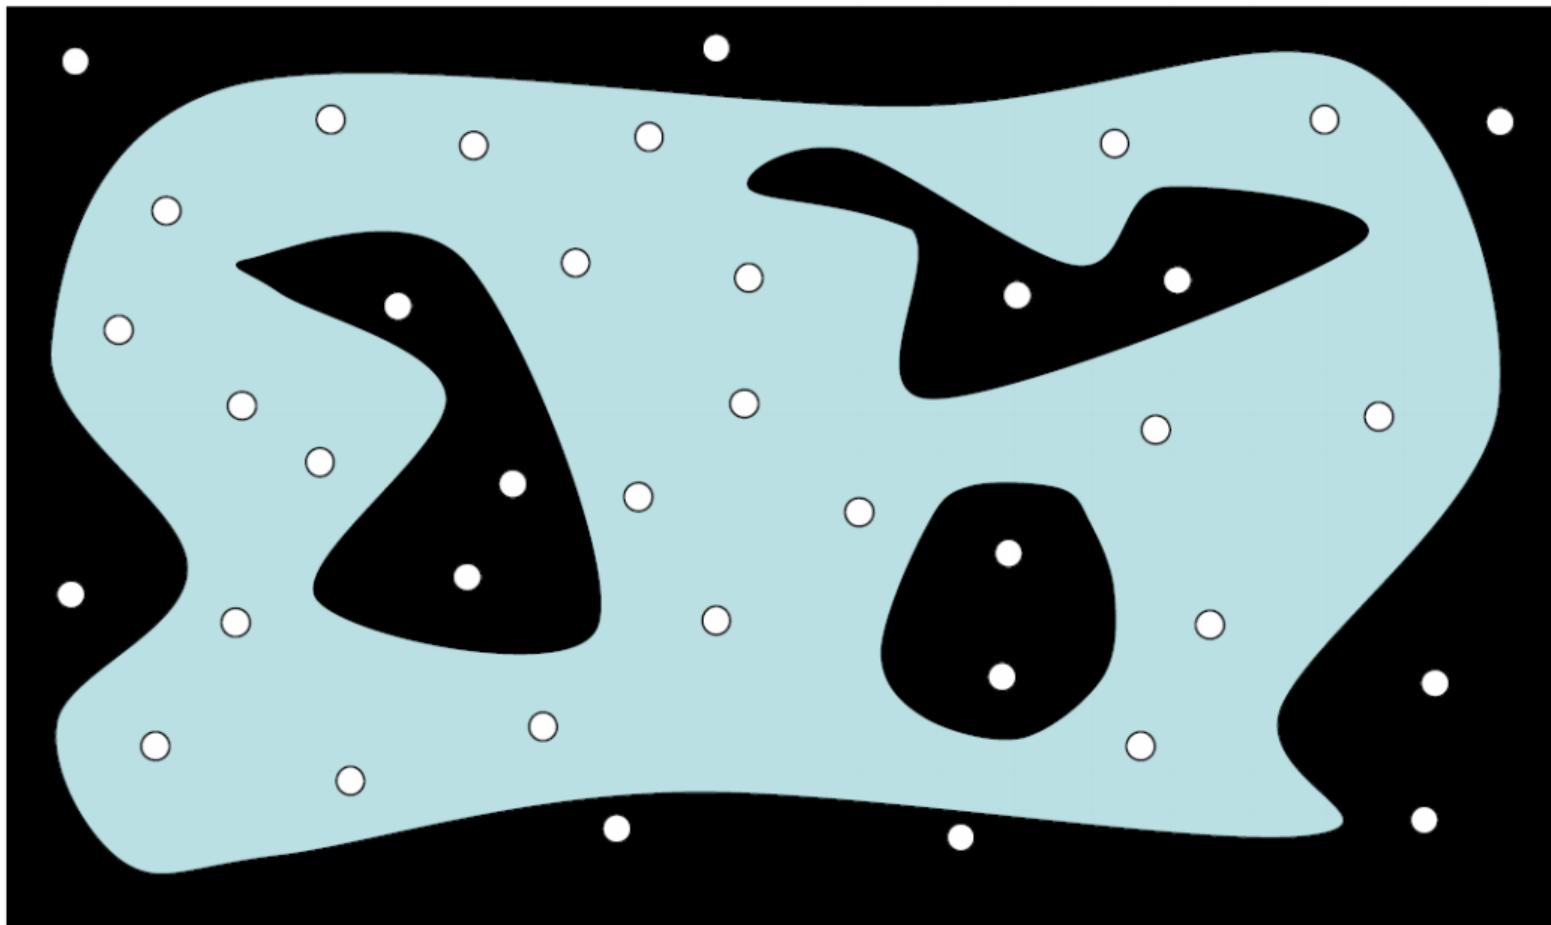
Potential Fields: Strengths/Weaknesses



Probabilistic Road Map (PRM)

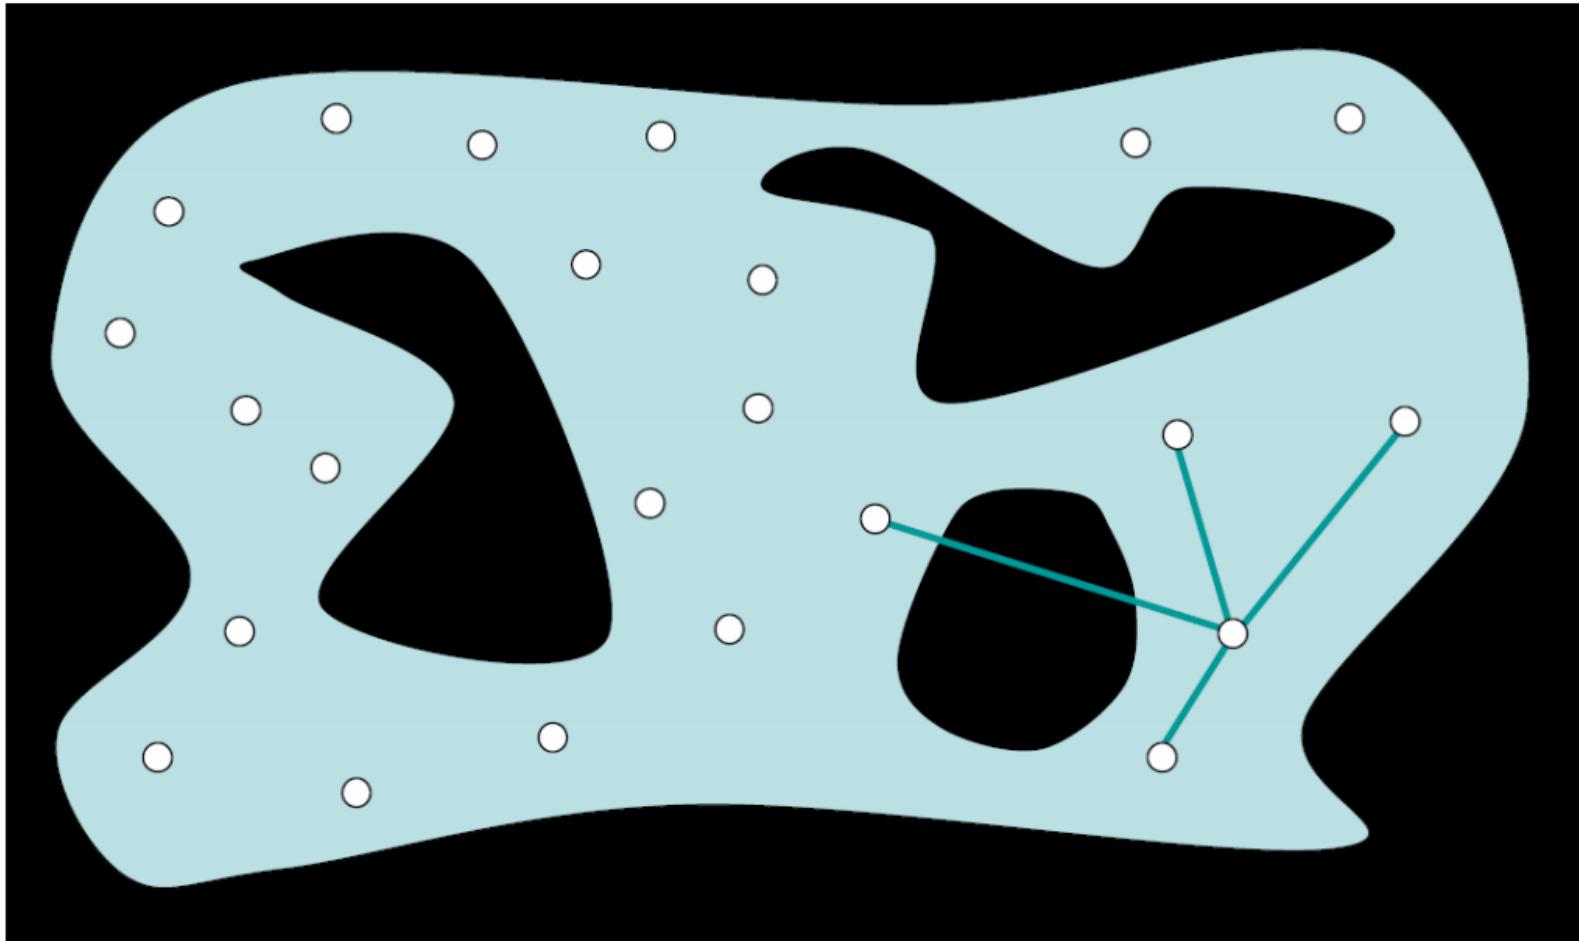


Probabilistic Road Map (PRM)



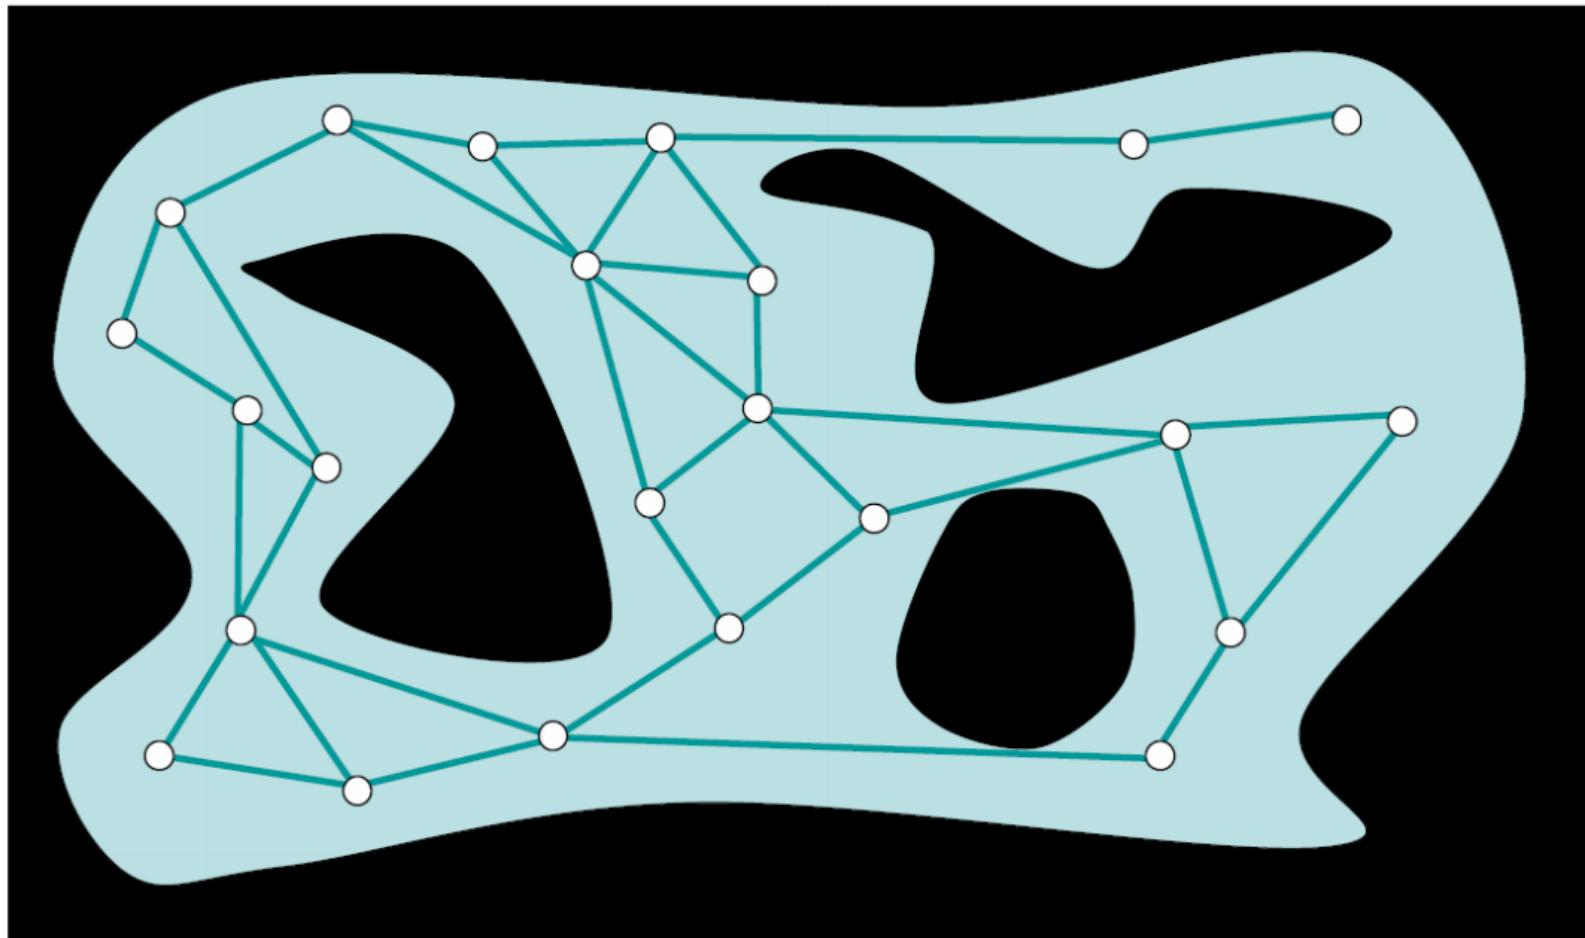
Sample random locations!

Probabilistic Road Map (PRM)



Remove points in forbidden areas
Link each point to its K nearest neighbors

Probabilistic Road Map (PRM)

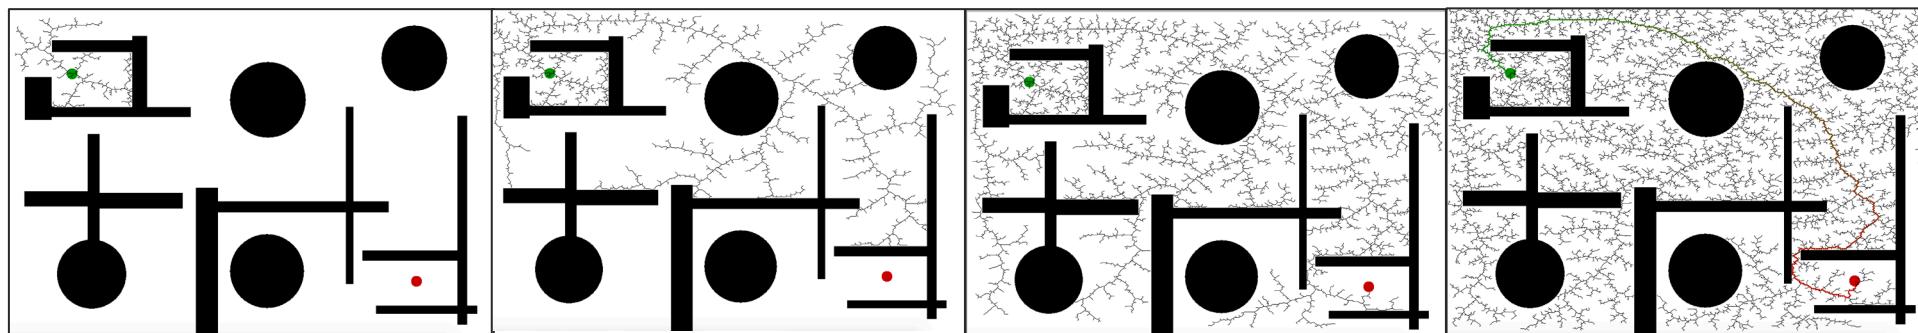


Remove edges crossing forbidden areas

How to sample points?

- Uniformly vs randomly
- Sample more near places with few neighbors
- Bias samples to exist near obstacles
- Use human-provided waypoints
- Something better?

Rapidly-exploring Random Trees (RRT)



Rapidly Exploring Random Trees [RRT]

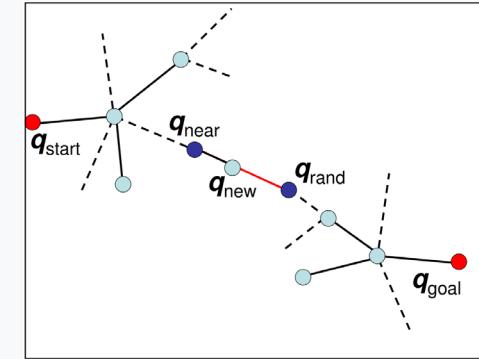
Basic idea:

- Build up a tree through generating “next states” in the tree by executing **random controls**
- However: to ensure good coverage,
we need something better than
the above!!

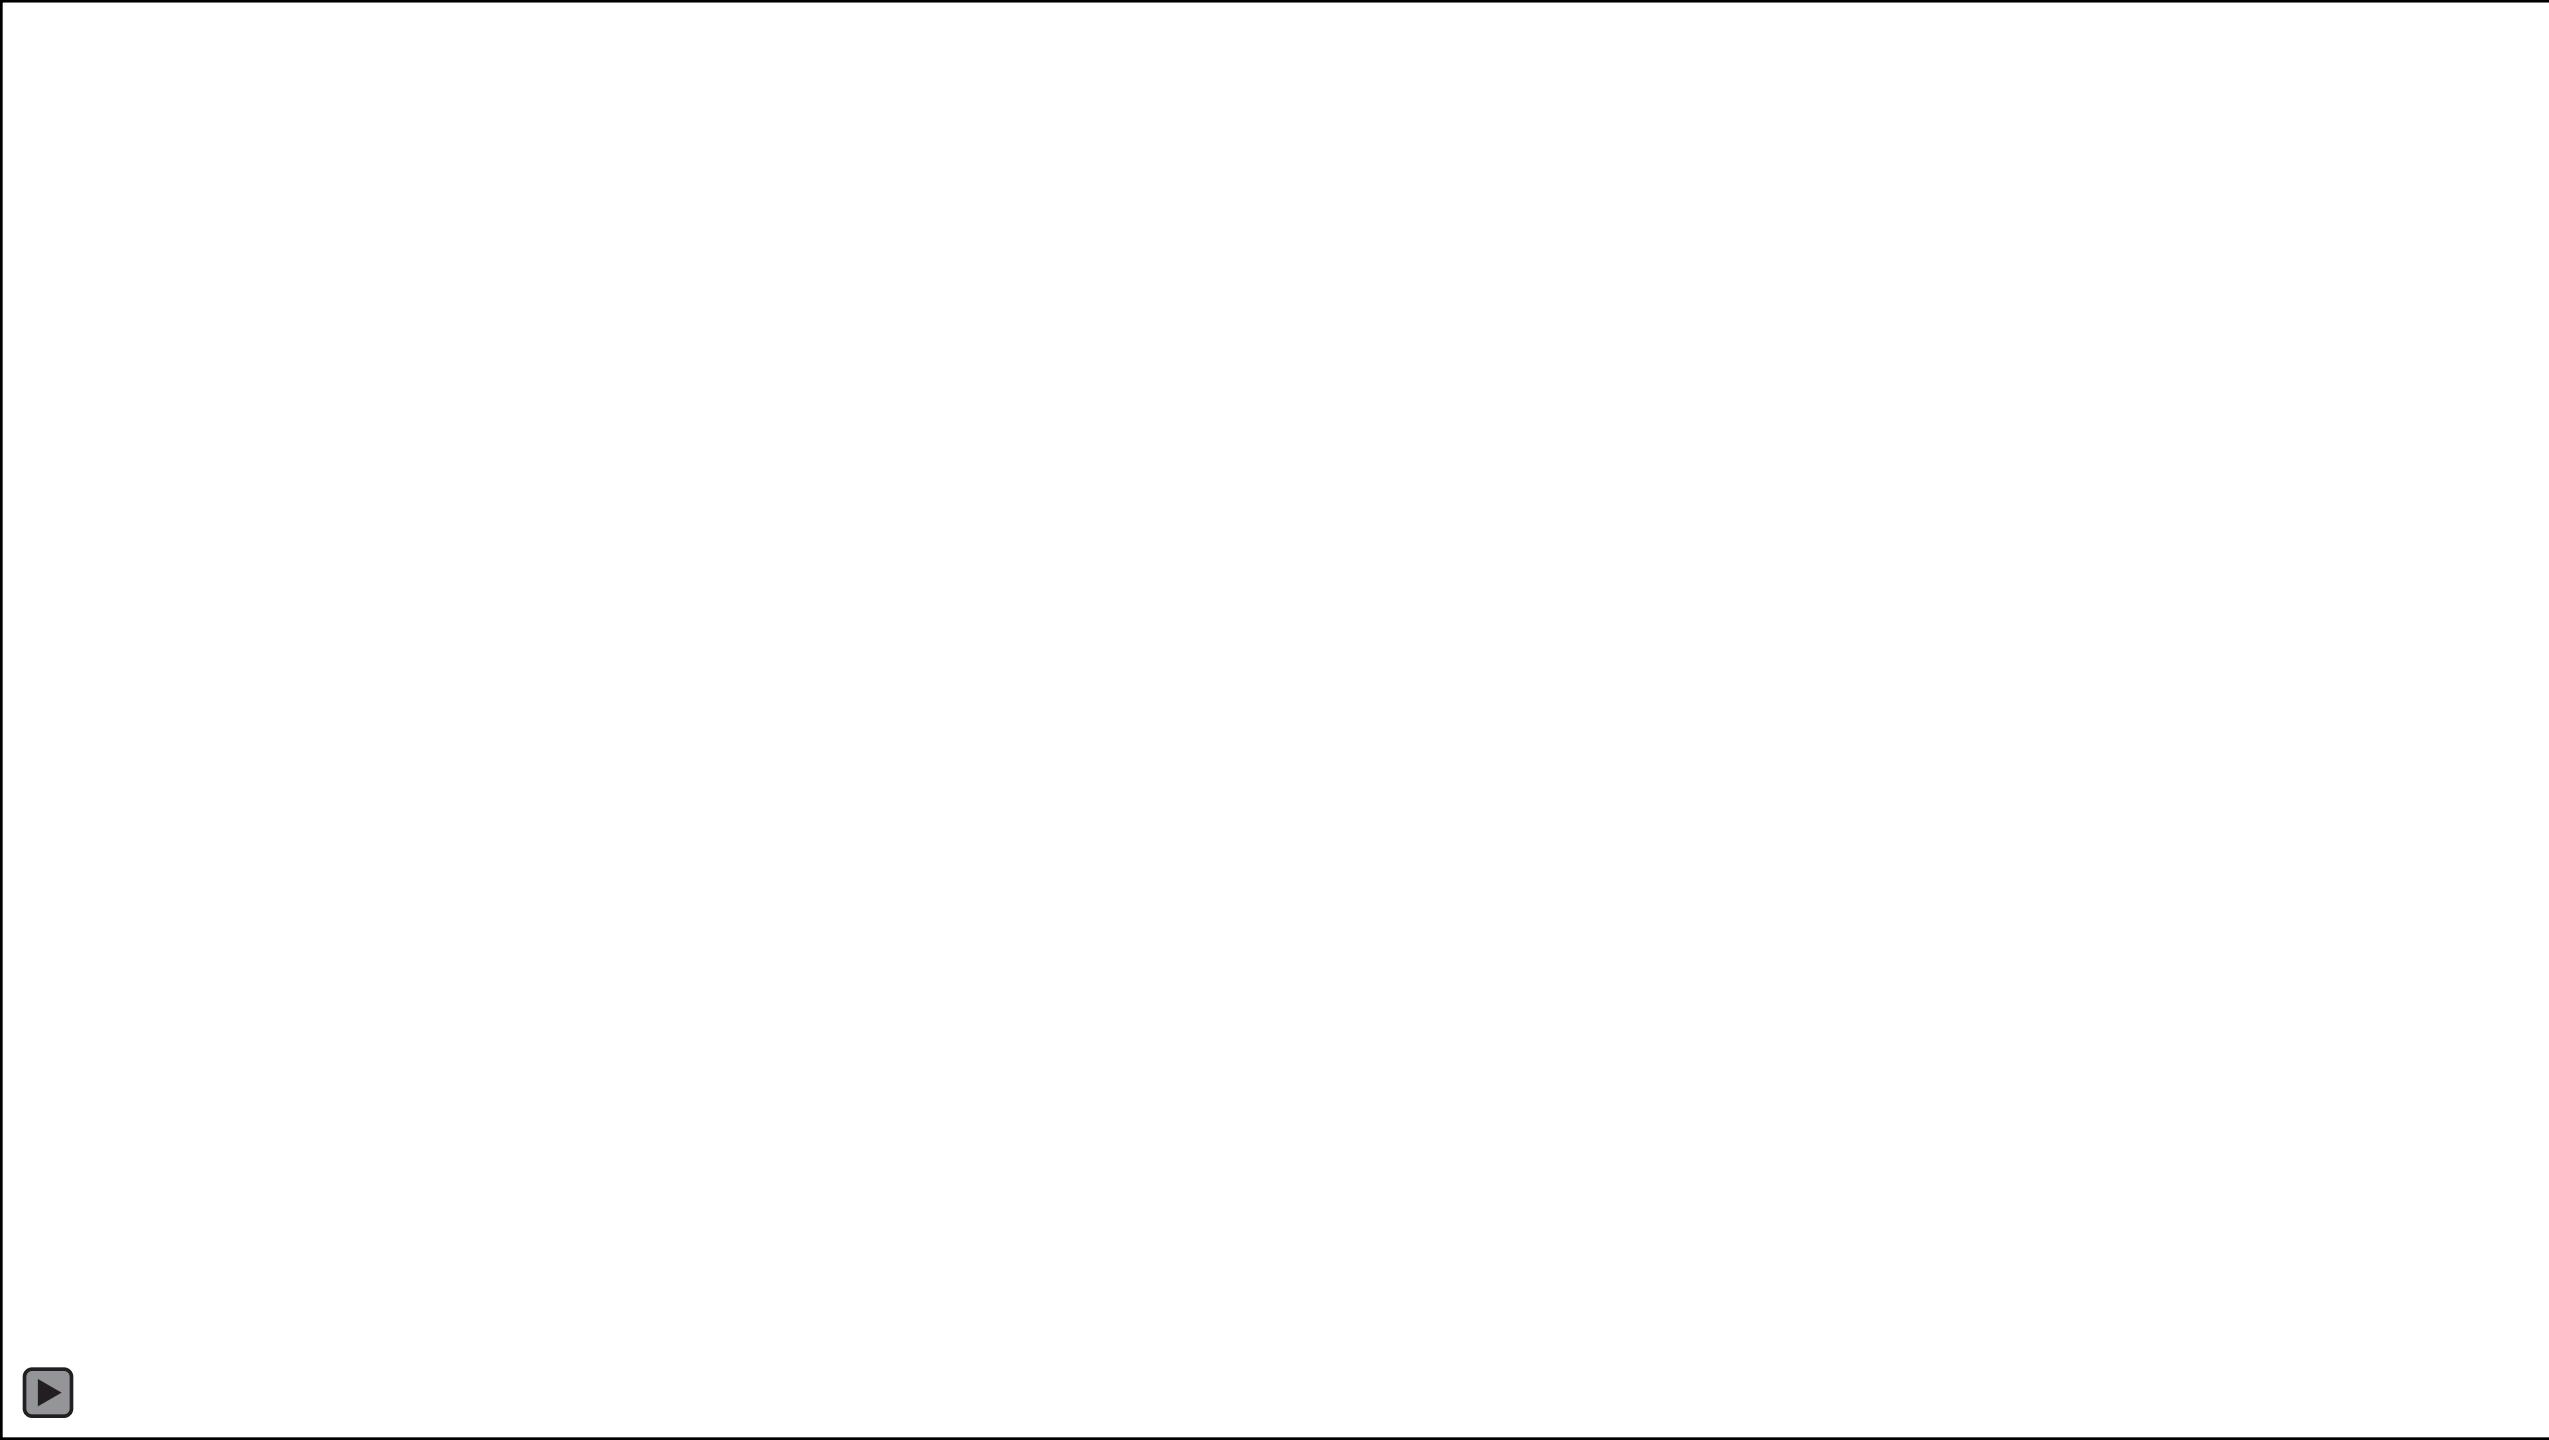
Rapidly-exploring Random Trees

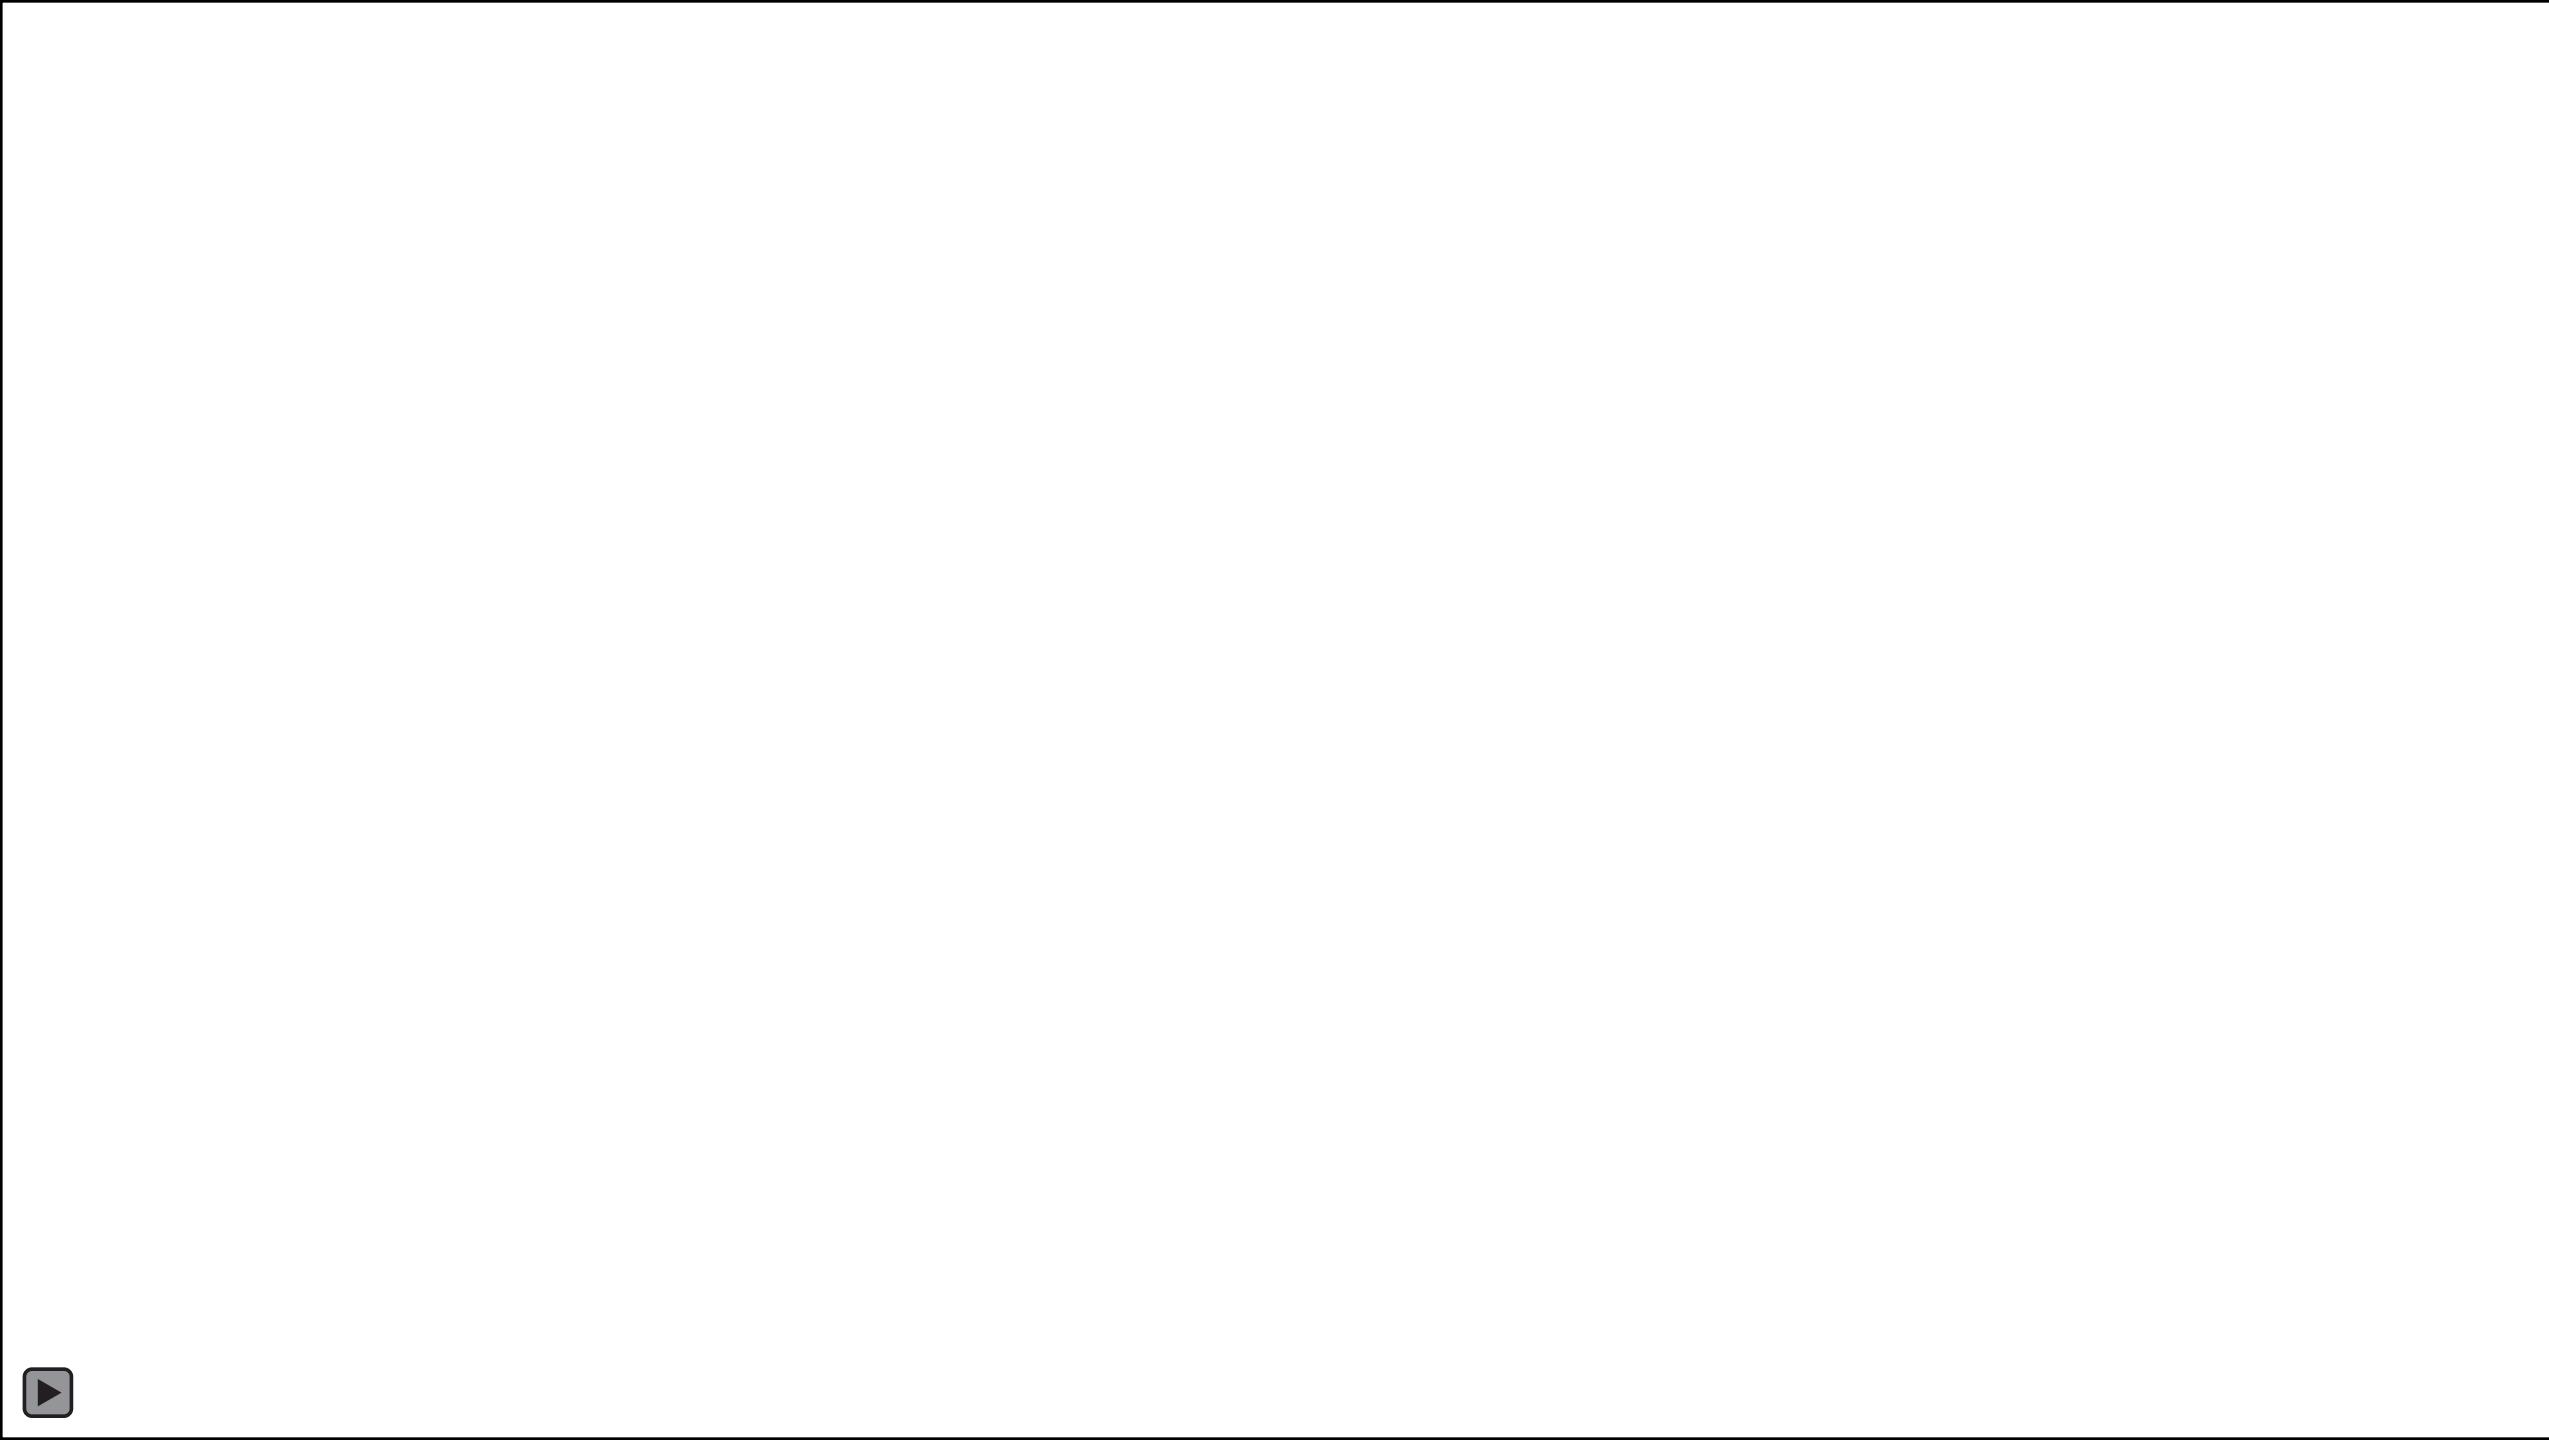
```
Algorithm BuildRRT
    Input: Initial configuration  $q_{init}$ , number of vertices in RRT  $K$ , incremental distance  $\Delta q$ )
    Output: RRT graph  $G$ 

     $G.init(q_{init})$ 
    for  $k = 1$  to  $K$ 
         $q_{rand} \leftarrow \text{RAND\_CONF}()$ 
         $q_{near} \leftarrow \text{NEAREST\_VERTEX}(q_{rand}, G)$ 
         $q_{new} \leftarrow \text{NEW\_CONF}(q_{near}, q_{rand}, \Delta q)$ 
         $G.add\_vertex(q_{new})$ 
         $G.add\_edge(q_{near}, q_{new})$ 
    return  $G$ 
```



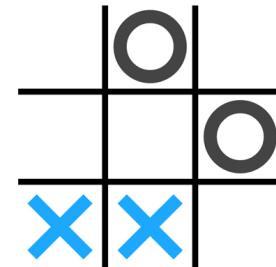
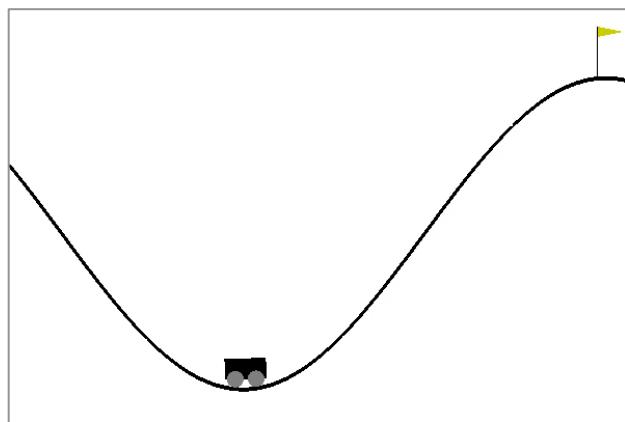
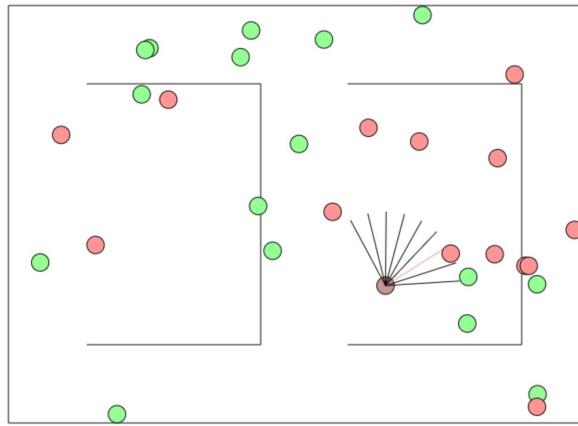
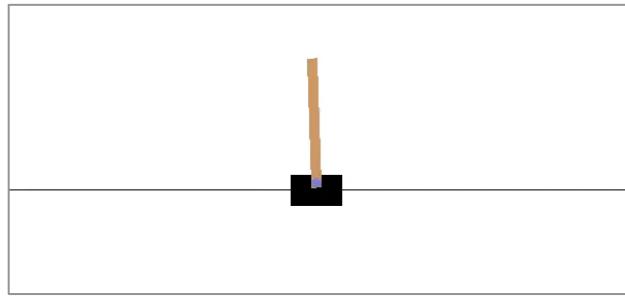
How can we make use of these representations?
Search algorithms provide a way to find a path!





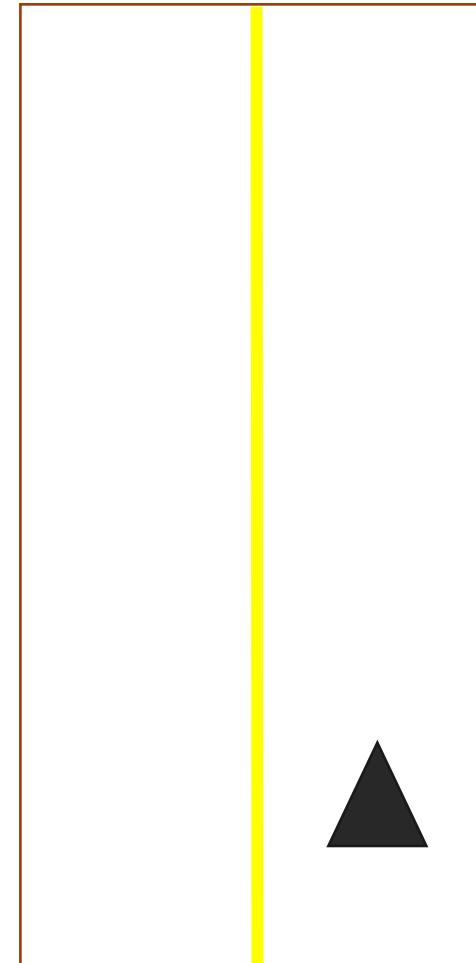
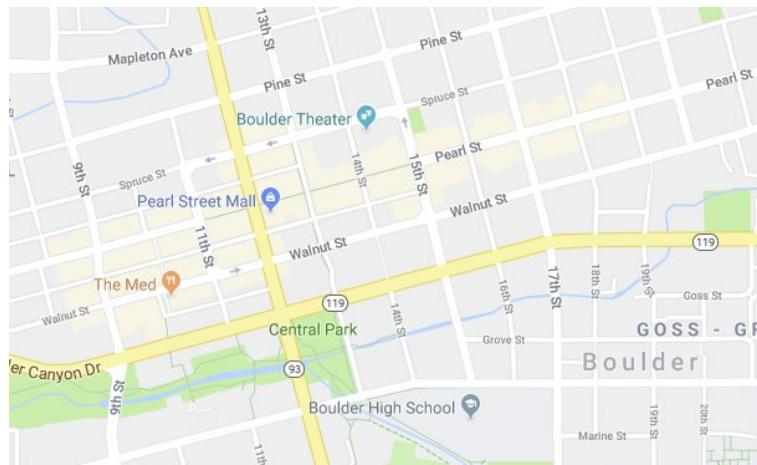


State Representation is Critical



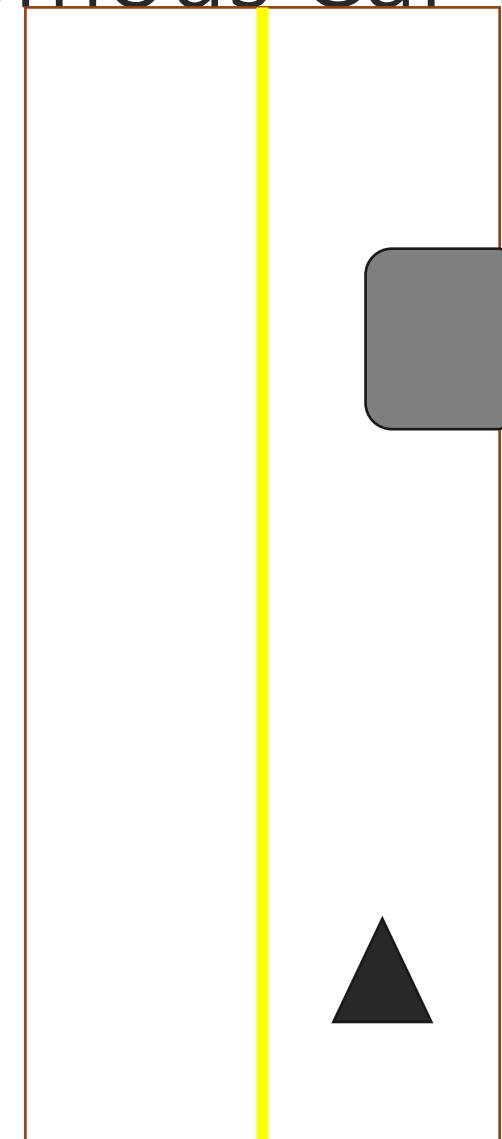
Exercise: Building an Autonomous Car

1. World Representation
2. Planning System

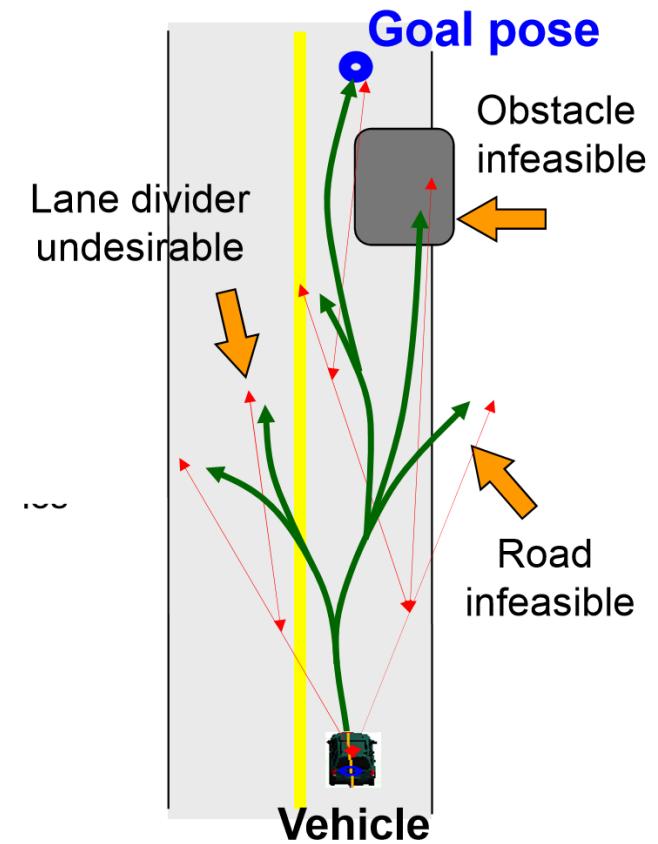


Exercise: Building an Autonomous Car

How can we get our
agent to accommodate
obstacles?



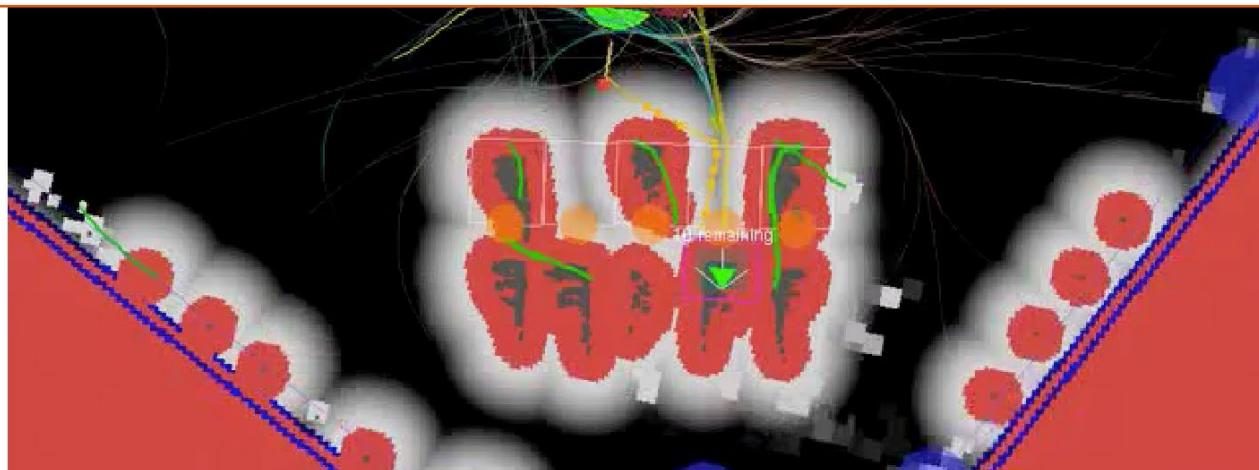
Exercise: Building an Autonomous Car



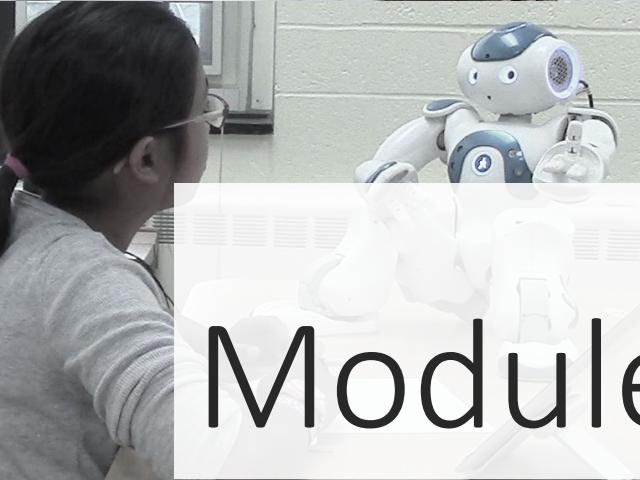
RRT for Autonomous Parking



How can we make use of these representations?
Search algorithms provide a way to find a path!



Chapter 13



Module 2 – COMPUTATION

Part III – Search Algorithms



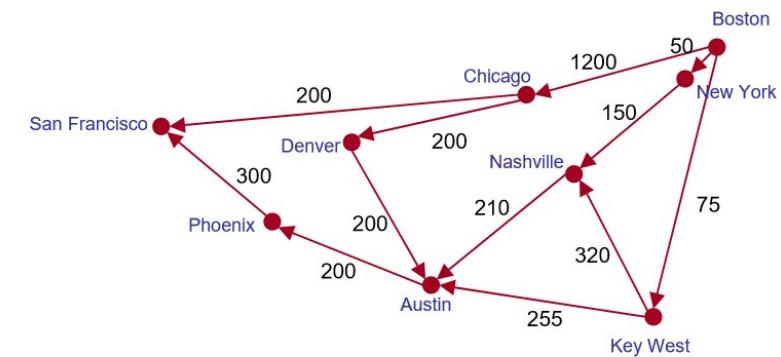
Search Algorithms

Allow us to **use** the world representations:

- Basics
- Uniform Cost
- Informed Search

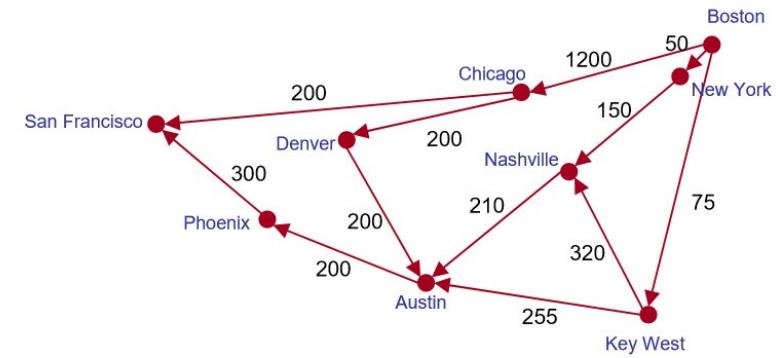
Goal Formulation

- Well defined function that identifies both the **goal states** and the **conditions** under which to achieve it
 - E.g. “Fly from Boston to San Francisco”
- **Metrics** for optimality may depend on:
 - Least amount of money
 - Fewest number of transfers
 - Shortest amount of time in the air
 - Shortest amount of time in airports

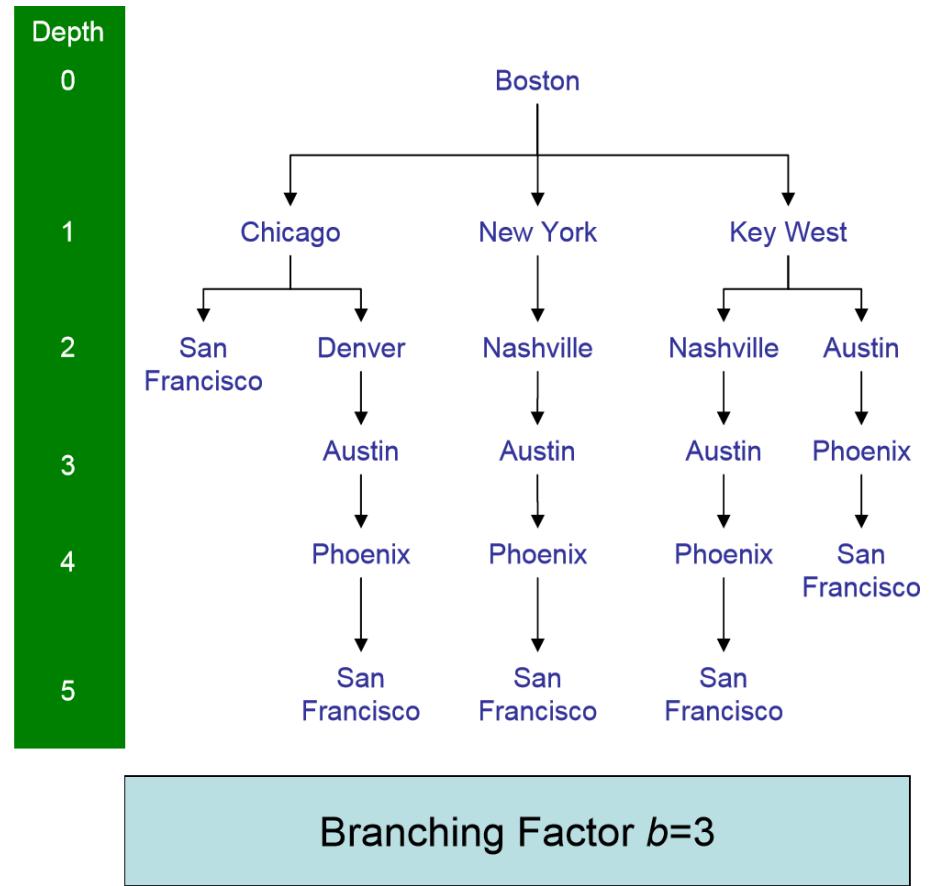
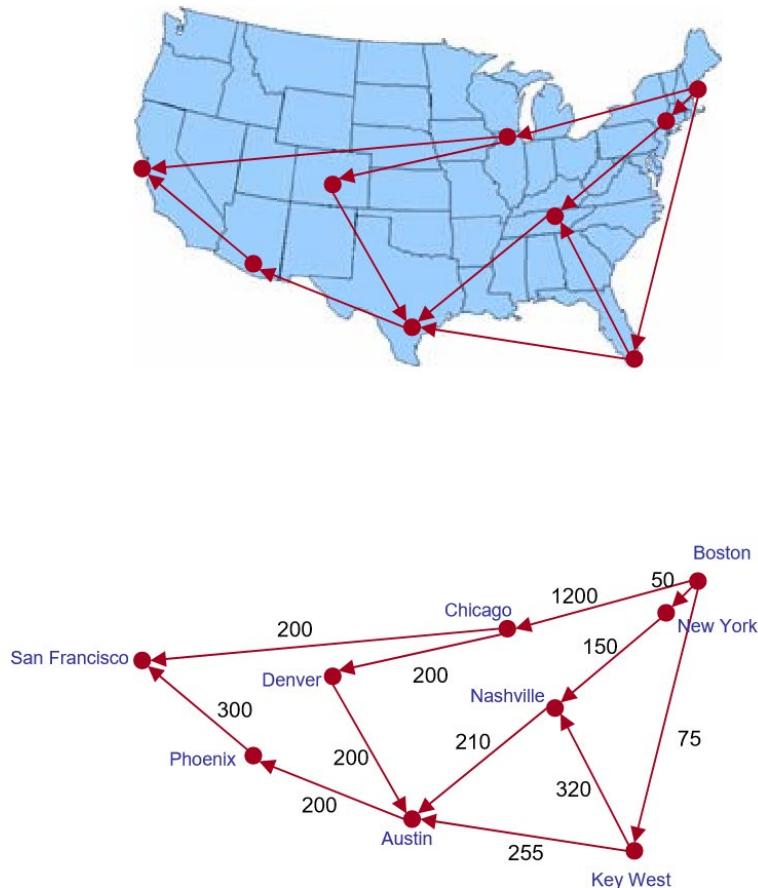


Problem formulation

- **Well defined** problems
 - Fully observable
 - Deterministic
 - Discrete Set of possible actions
- **State Space:** the set of all states that are reachable from an initial state by any sequence of actions
- **Path:** sequence of actions leading from one state to another



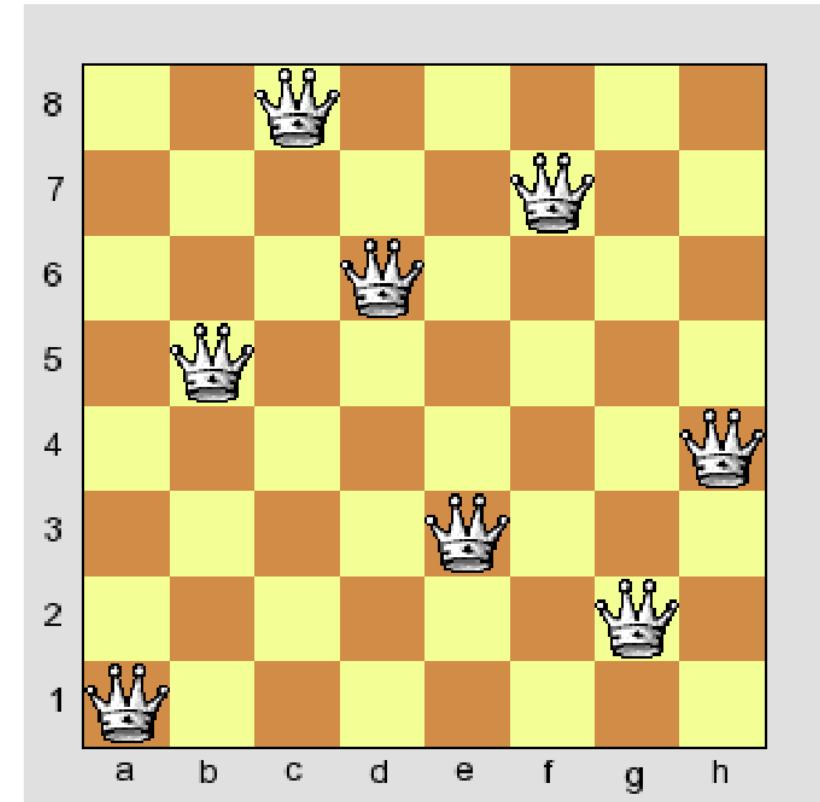
Search as a problem solving technique!!



Problem Formulation Matters!

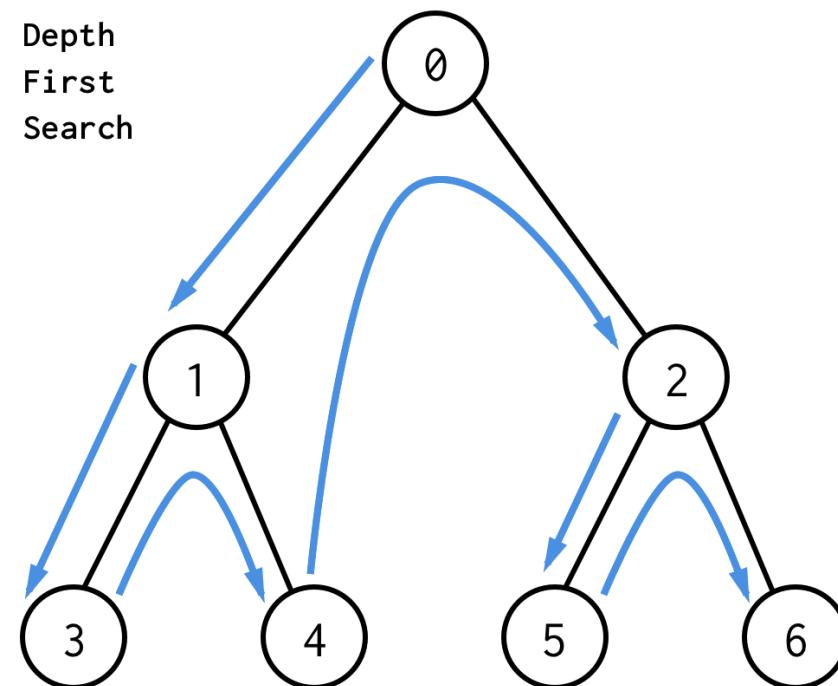
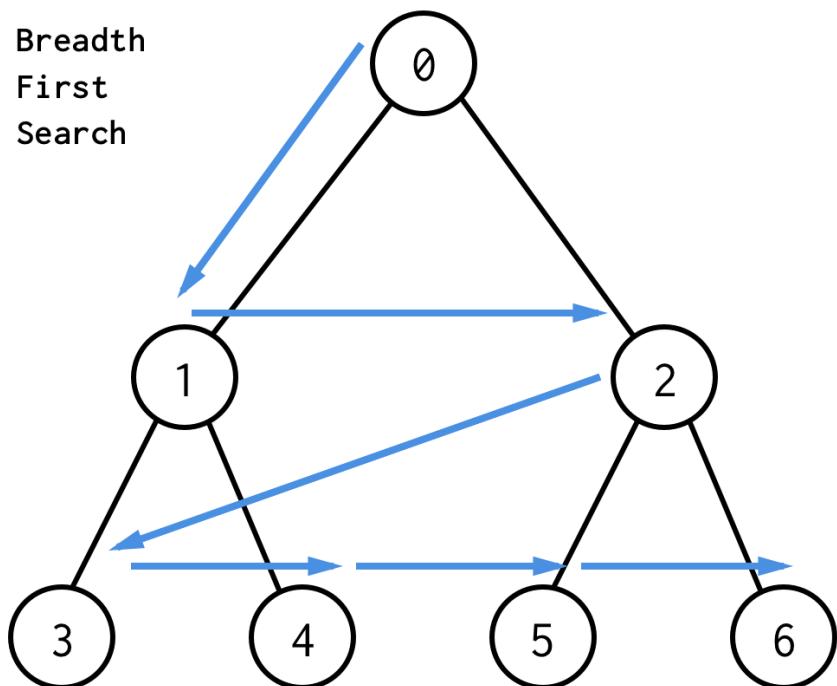
The 8 Queens Problem

- Formulation #1:
 - Place a queen on any open square
 - Repeat until all queens are placed
 - State space of $\frac{64!}{56!} = 1.78 * 10^{14}$
- Formulation #2:
 - Place a queen on any row 1 square
 - Place a queen on any row 2 square
 - State space of $8^8 = 1.68 * 10^7$
- Formulation #3:
 - Place a queen on any row 1 square
 - Place a queen on any row 2 square not sharing a column...
 - State space of $8! = 40,320$



Brute force becomes intractable for problems of $n \geq 20$, as $20! = 2.433 \times 10^{18}$

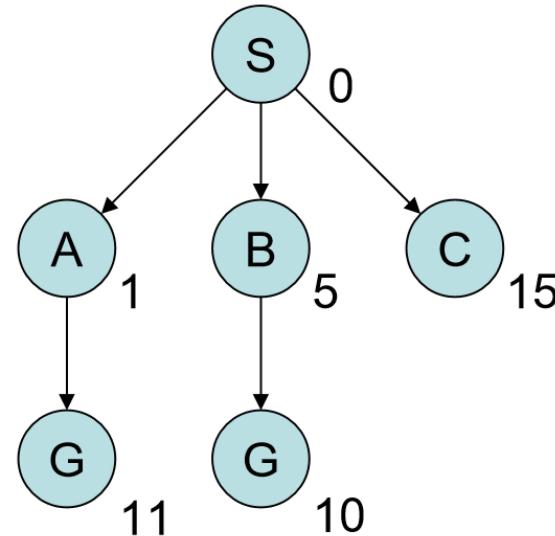
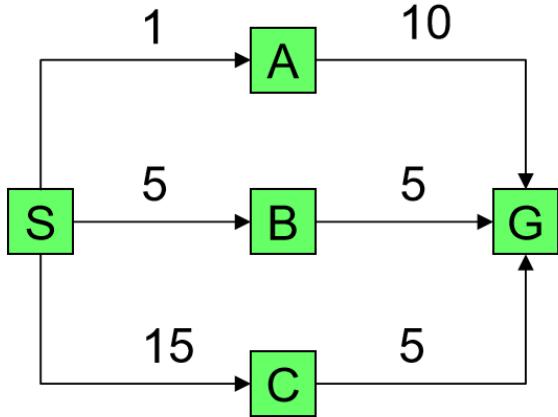
Breadth First vs Depth First search



- Finds the most **shallow** solution

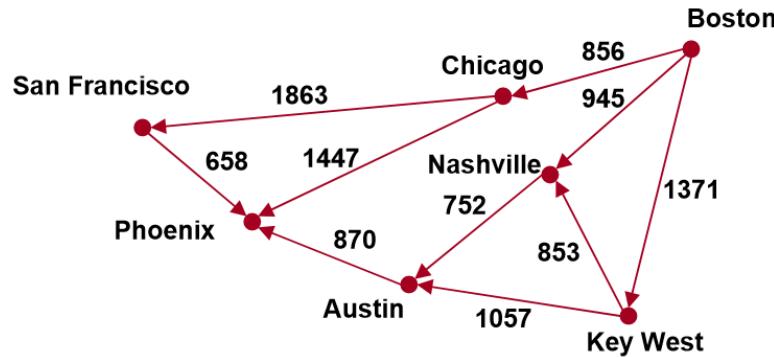
- Finds the most **deep** solution

Uniform Cost Search



- Travel from the start (S) to the goal (G)
- Cost associated with each link
- Always expand the fringe node with the lowest cost
- Breadth-first search is uniform search with $\text{cost}=\text{depth}$

Greedy Best-First-Search



- Minimize estimated cost to reach a goal (in this case, the distance to Phoenix)

	Straight Line Distance to Phoenix
Boston	2299
Chicago	1447
Nashville	1444
Key West	1927
Austin	870
San Francisco	658

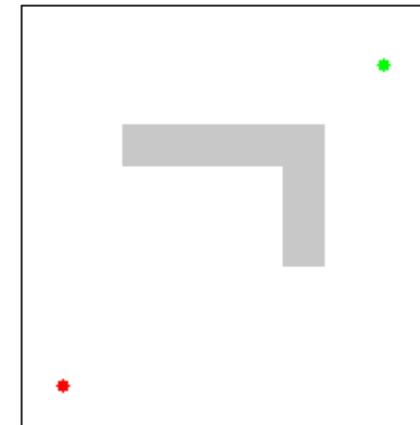
Search Algorithms: Uniform Cost

Dijkstra's Algorithm:

```
1 function Dijkstra(Graph, source):
2
3     create vertex set Q
4
5     for each vertex v in Graph:           // Initialization
6         dist[v] ← INFINITY               // Unknown distance from source to v
7         prev[v] ← UNDEFINED              // Previous node in optimal path from source
8         add v to Q                      // All nodes initially in Q (unvisited nodes)
9
10    dist[source] ← 0                   // Distance from source to source
11
12    while Q is not empty:
13        u ← vertex in Q with min dist[u] // Node with the least distance
14                                // will be selected first
15
16        remove u from Q
17
18        for each neighbor v of u:          // where v is still in Q.
19            alt ← dist[u] + length(u, v)
20            if alt < dist[v]:              // A shorter path to v has been found
21                dist[v] ← alt
22                prev[v] ← u
23
24    return dist[], prev[]
```

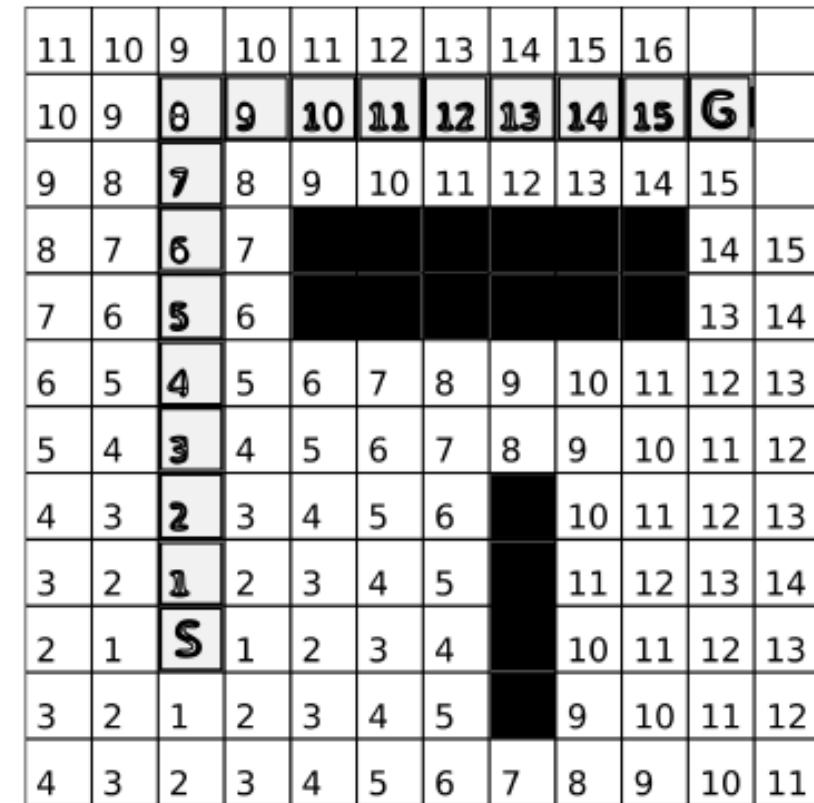
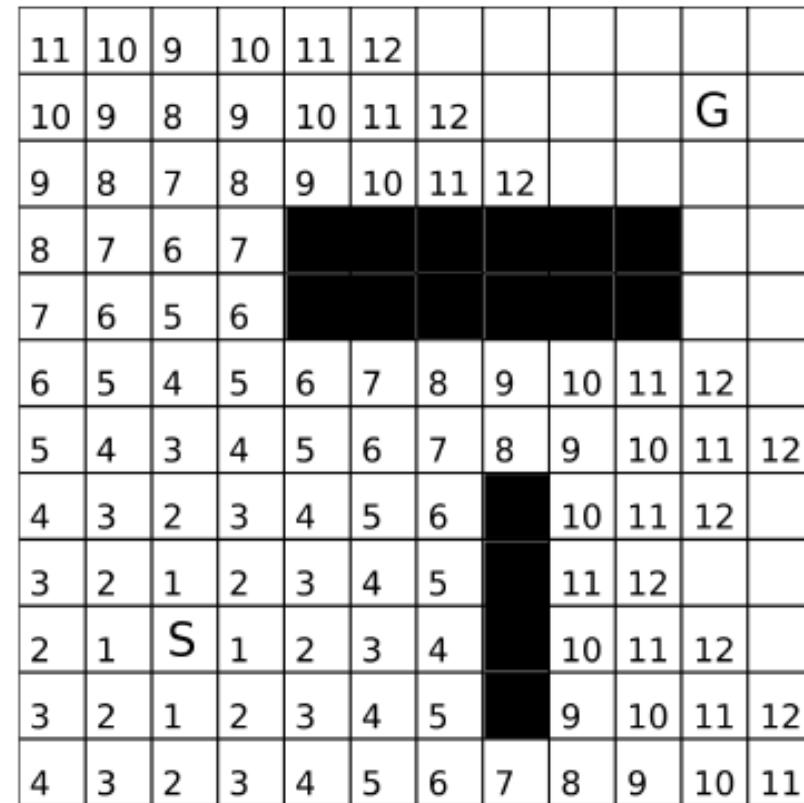
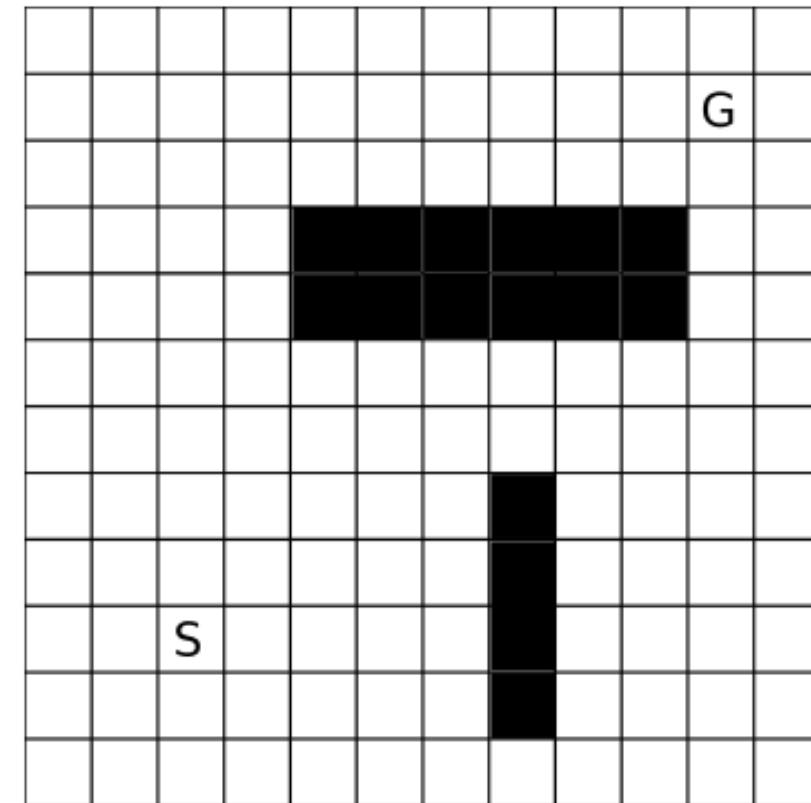
```
1 S ← empty sequence
2 u ← target
3 if prev[u] is defined or u = source:
4     while u is defined:
5         insert u at the beginning of S
6         u ← prev[u]
7
8     // Do something only if the vertex is reachable
9     // Construct the shortest path with a stack S
10    // Push the vertex onto the stack
11    // Traverse from target to source
```

Finds shortest paths
between nodes in a graph



Dijkstra's Algorithm on a Grid

Problem: A lot of useless exploration



Informed Search

- Dijkstra's Algorithm:
- Lots of wasted exploration!
- We usually know which two nodes we want to find a path between:
 - Why compute all distances from source?
- **Heuristic**: A function that ranks choices in a search algorithm to help choose how to branch/explore.

Finds shortest path
between two nodes in
a graph



A* Search Algorithm

Finds shortest path
between two nodes in
a graph

- Adds heuristic to Dijkstra's Algorithm

to **bias exploration!**

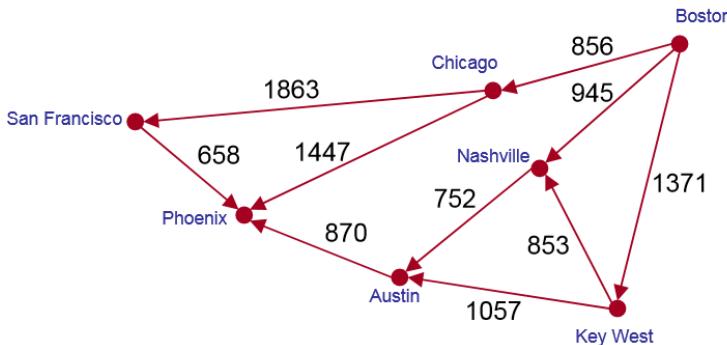
- Insight:

$$f(n) = g(n) + h(n)$$

- $f(n)$ = Cost so far + Est. cost to go

Admissible Heuristic: An **underestimate** of the shortest possible path from node n to the goal.

A* Search Algorithm



- Combine Greedy search with Uniform Cost Search
- Minimize the total path cost (f) =
actual path so far (g) +
estimate of future path to goal (h)

	Distance to Phoenix
Boston	2299
Chicago	1447
Nashville	1444
Key West	1927
Austin	870
San Francisco	658

A* Search Algorithm

- Complete: Yes

(If an answer exists, this will find it)

- Optimal: Yes

(Returned answer will be the best possible)

- Optimally Efficient: Yes

(No algorithm with same heuristic will expand fewer nodes)

- Finding good admissible heuristics is challenging!

```
function reconstruct_path(cameFrom, current)
    total_path := [current]
    while current in cameFrom.Keys:
        current := cameFrom[current]
        total_path.append(current)
    return total_path
```

```
function A*(start, goal)
    // The set of nodes already evaluated
    closedSet := {}

    // The set of currently discovered nodes that are not evaluated yet.
    // Initially, only the start node is known.
    openSet := {start}

    // For each node, which node it can most efficiently be reached from.
    // If a node can be reached from many nodes, cameFrom will eventually contain the
    // most efficient previous step.
    cameFrom := the empty map

    // For each node, the cost of getting from the start node to that node.
    gScore := map with default value of Infinity
    gScore[start] := 0

    // For each node, the total cost of getting from the start node to the goal
    // by passing by that node. That value is partly known, partly heuristic.
    fScore := map with default value of Infinity
    fScore[start] := heuristic_cost_estimate(start, goal)

    while openSet is not empty
        current := the node in openSet having the lowest fScore[] value
        if current = goal
            return reconstruct_path(cameFrom, current)

        openSet.Remove(current)
        closedSet.Add(current)

        for each neighbor of current
            if neighbor in closedSet
                continue      // Ignore the neighbor which is already evaluated.

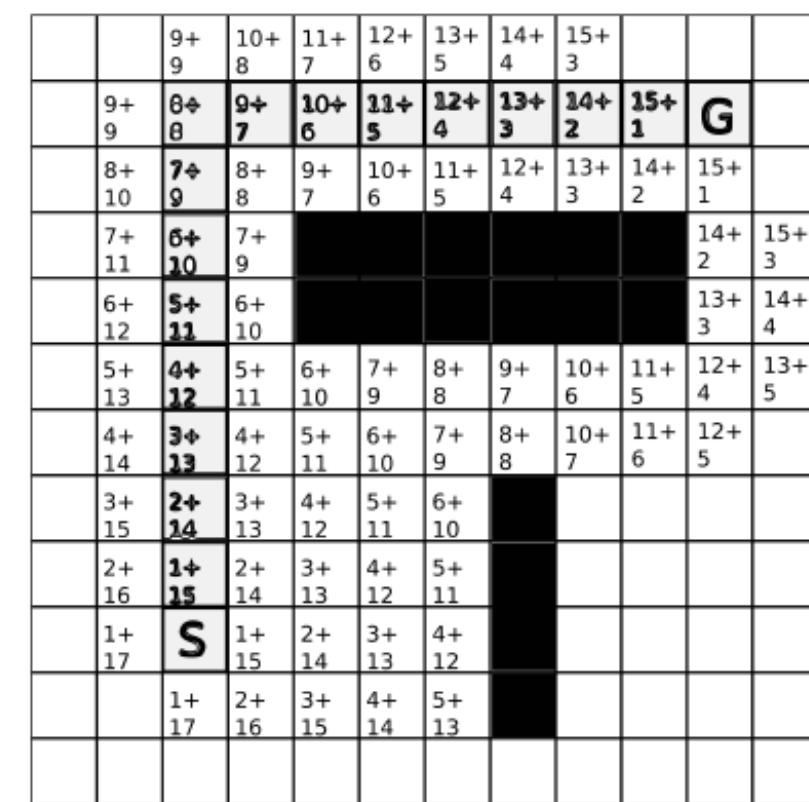
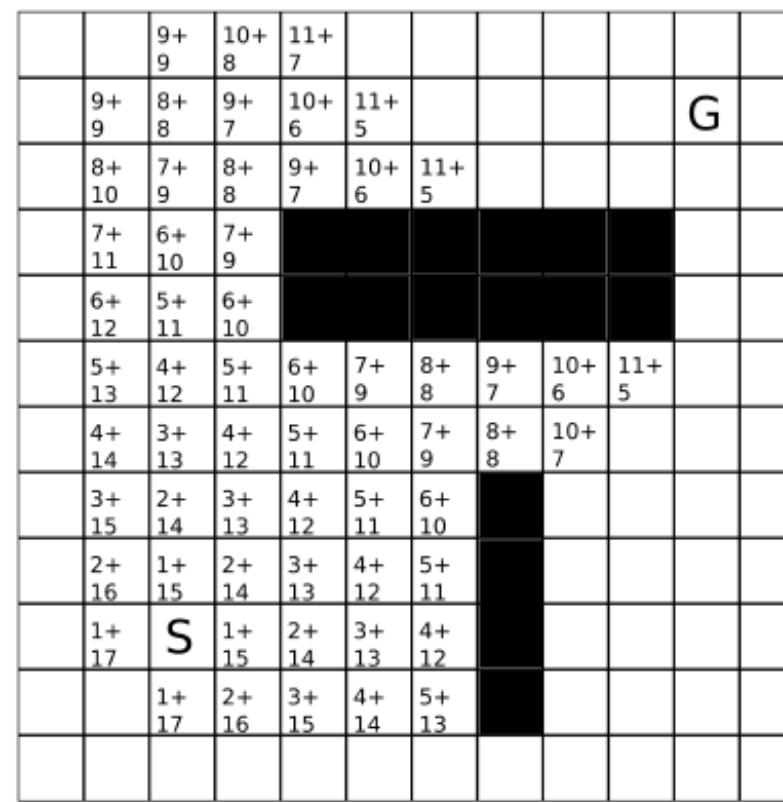
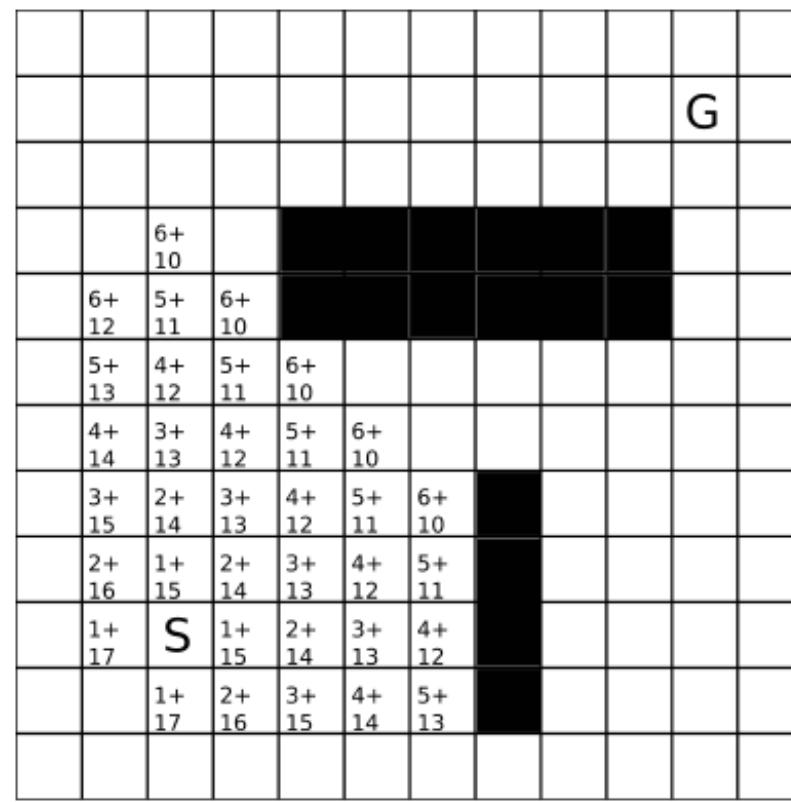
            if neighbor not in openSet // Discover a new node
                openSet.Add(neighbor)

            // The distance from start to a neighbor
            tentative_gScore := gScore[current] + dist_between(current, neighbor)
            if tentative_gScore >= gScore[neighbor]
                continue      // This is not a better path.

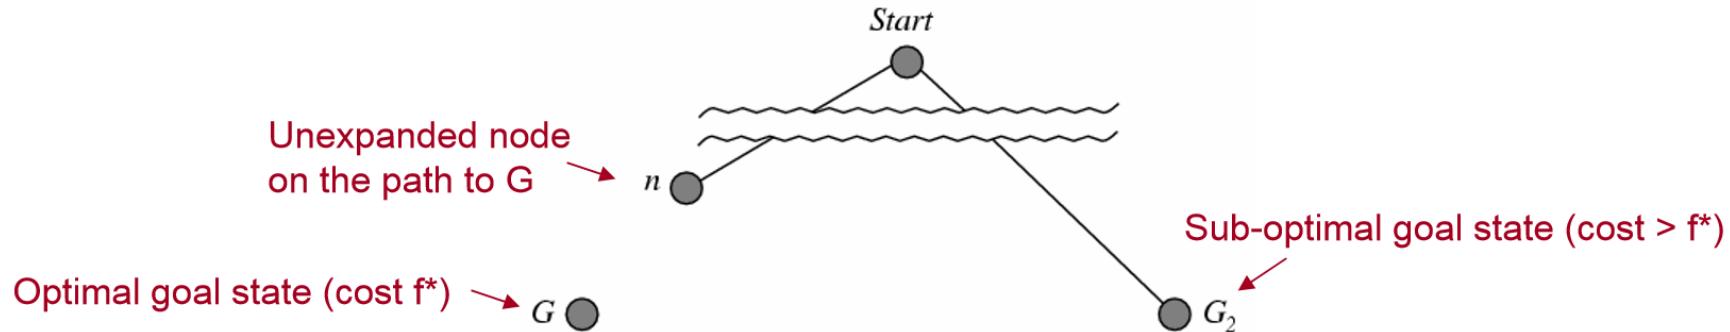
            // This path is the best until now. Record it!
            cameFrom[neighbor] := current
            gScore[neighbor] := tentative_gScore
            fScore[neighbor] := gScore[neighbor] + heuristic_cost_estimate(neighbor, goal)

    return failure
```

Dijkstra plus directional heuristic: A*



Optimality of A*

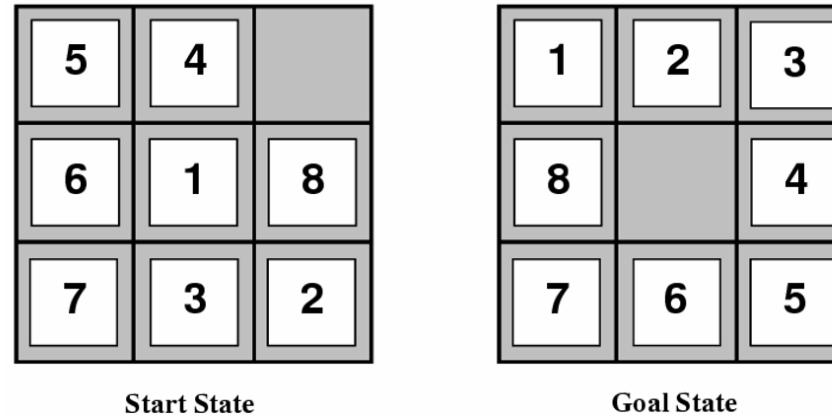


- Assume that G_2 has been chosen for expansion over n
- Because h is admissible
$$f^* \geq f(n)$$
- If n is not chosen for expansion over G_2 , we must have
$$f(n) \geq f(G_2)$$
- Combining these, we get
$$f^* \geq f(G_2)$$
- However, this violates our assertion that G_2 is sub-optimal
- Therefore, A* never selects a sub-optimal goal for expansion

Many variants of A* exist... (D* [lite])



Heuristic Selection



- Must be admissible (never over-estimate)
- Heuristics for the 8-Puzzle

Heuristic Selection Effects

d	Search Cost			Effective Branching Factor		
	IDS	A*(h_1)	A*(h_2)	IDS	A*(h_1)	A*(h_2)
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	364404	227	73	2.78	1.42	1.24
14	3473941	539	113	2.83	1.44	1.23
16	—	1301	211	—	1.45	1.25
18	—	3056	363	—	1.46	1.26
20	—	7276	676	—	1.47	1.27
22	—	18094	1219	—	1.48	1.28
24	—	39135	1641	—	1.48	1.26

- Compare iterative-deepening with A* using h_1 (# misplaced tiles) and h_2 (city block distance)
- Effective branching factor b^*
 - Number of expanded nodes = $1 + b^* + (b^*)^2 + \dots + (b^*)^{\text{depth}}$
 - b^* remains relatively constant across many measurements

How do humans solve the path/motion planning problem?

Take-home lessons

- First step in addressing a planning problem is choosing a suitable map representation
- Reduce robot to a point-mass by inflating obstacles
- Grid-based algorithms are complete, sampling-based ones probabilistically complete, but usually faster
- Most real planning problems require combination of multiple algorithms